

Extracting Clean Performance Models from Tainted Programs

Marcin Copik, Alexandru Calotoiu, Tobias Grosser, Nicolas Wicki, Felix Wolf, Torsten Hoefler


PPoPP 2021

3rd March 2021



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Performance Modeling: state of the art




Performance Modeling: state of the art




Performance Modeling: state of the art



Performance Modeling: state of the art



Performance Modeling: state of the art



Challenges in Automatic Performance Modeling

Parameters Identification



Select problem size s and ranks p as model parameters.

Challenges in Automatic Performance Modeling

Parameters Identification



Select problem size s and ranks p as model parameters.



Experiment Design

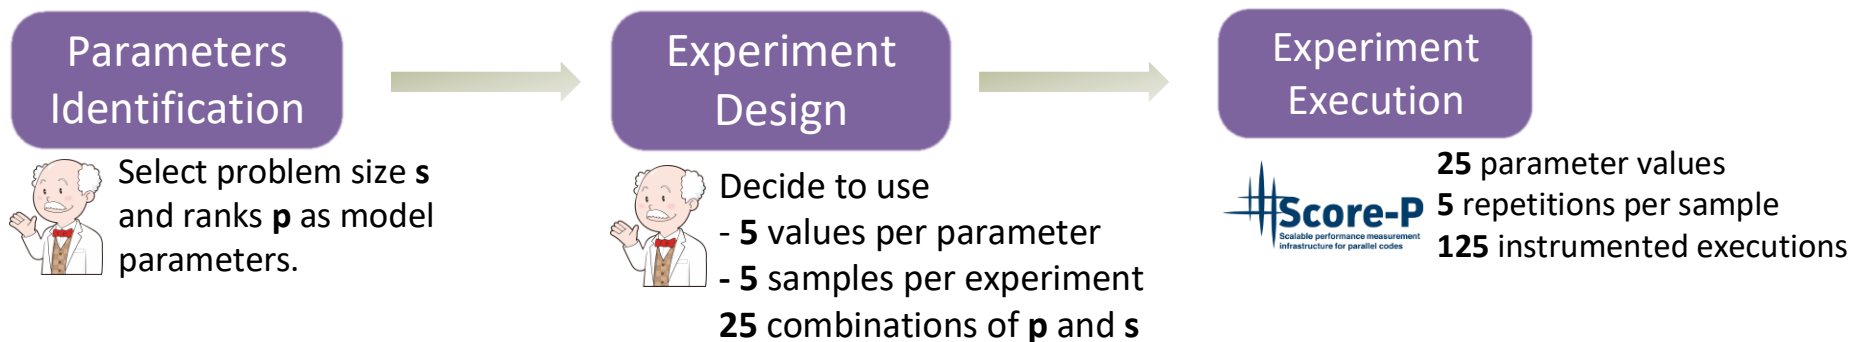


Decide to use


- 5 values per parameter
- 5 samples per experiment

25 combinations of p and s

Challenges in Automatic Performance Modeling



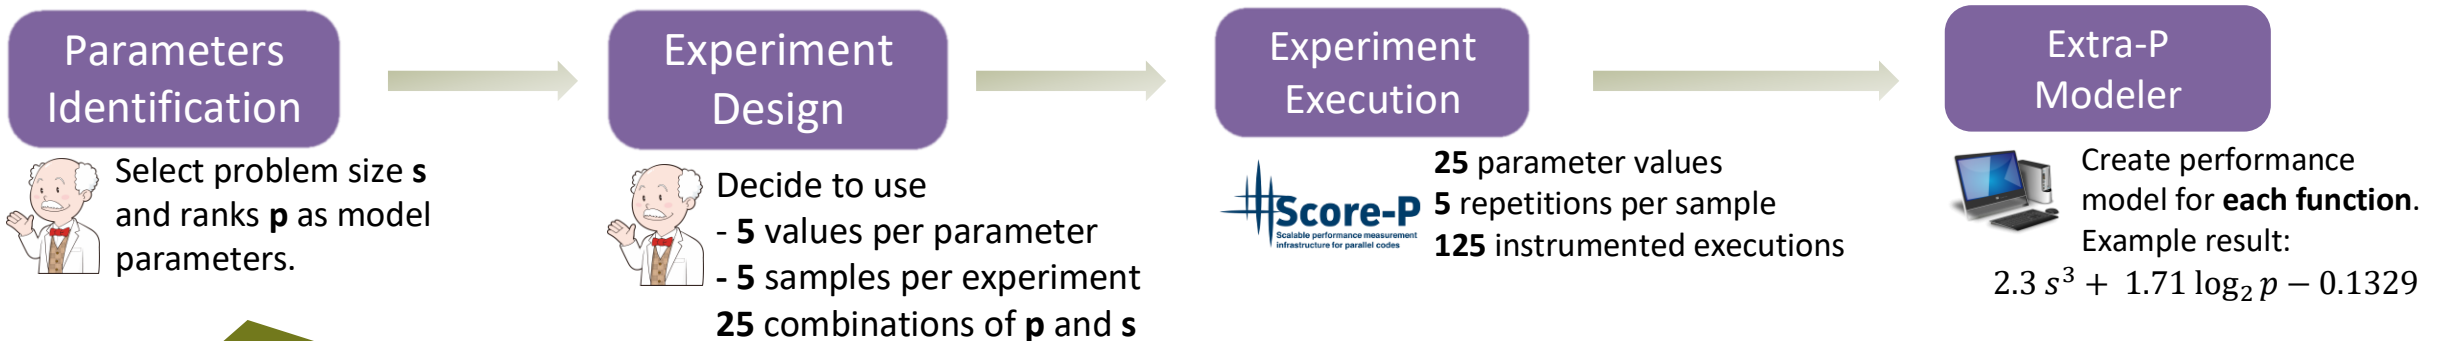
Challenges in Automatic Performance Modeling



Challenges in Automatic Performance Modeling



Challenges in Automatic Performance Modeling



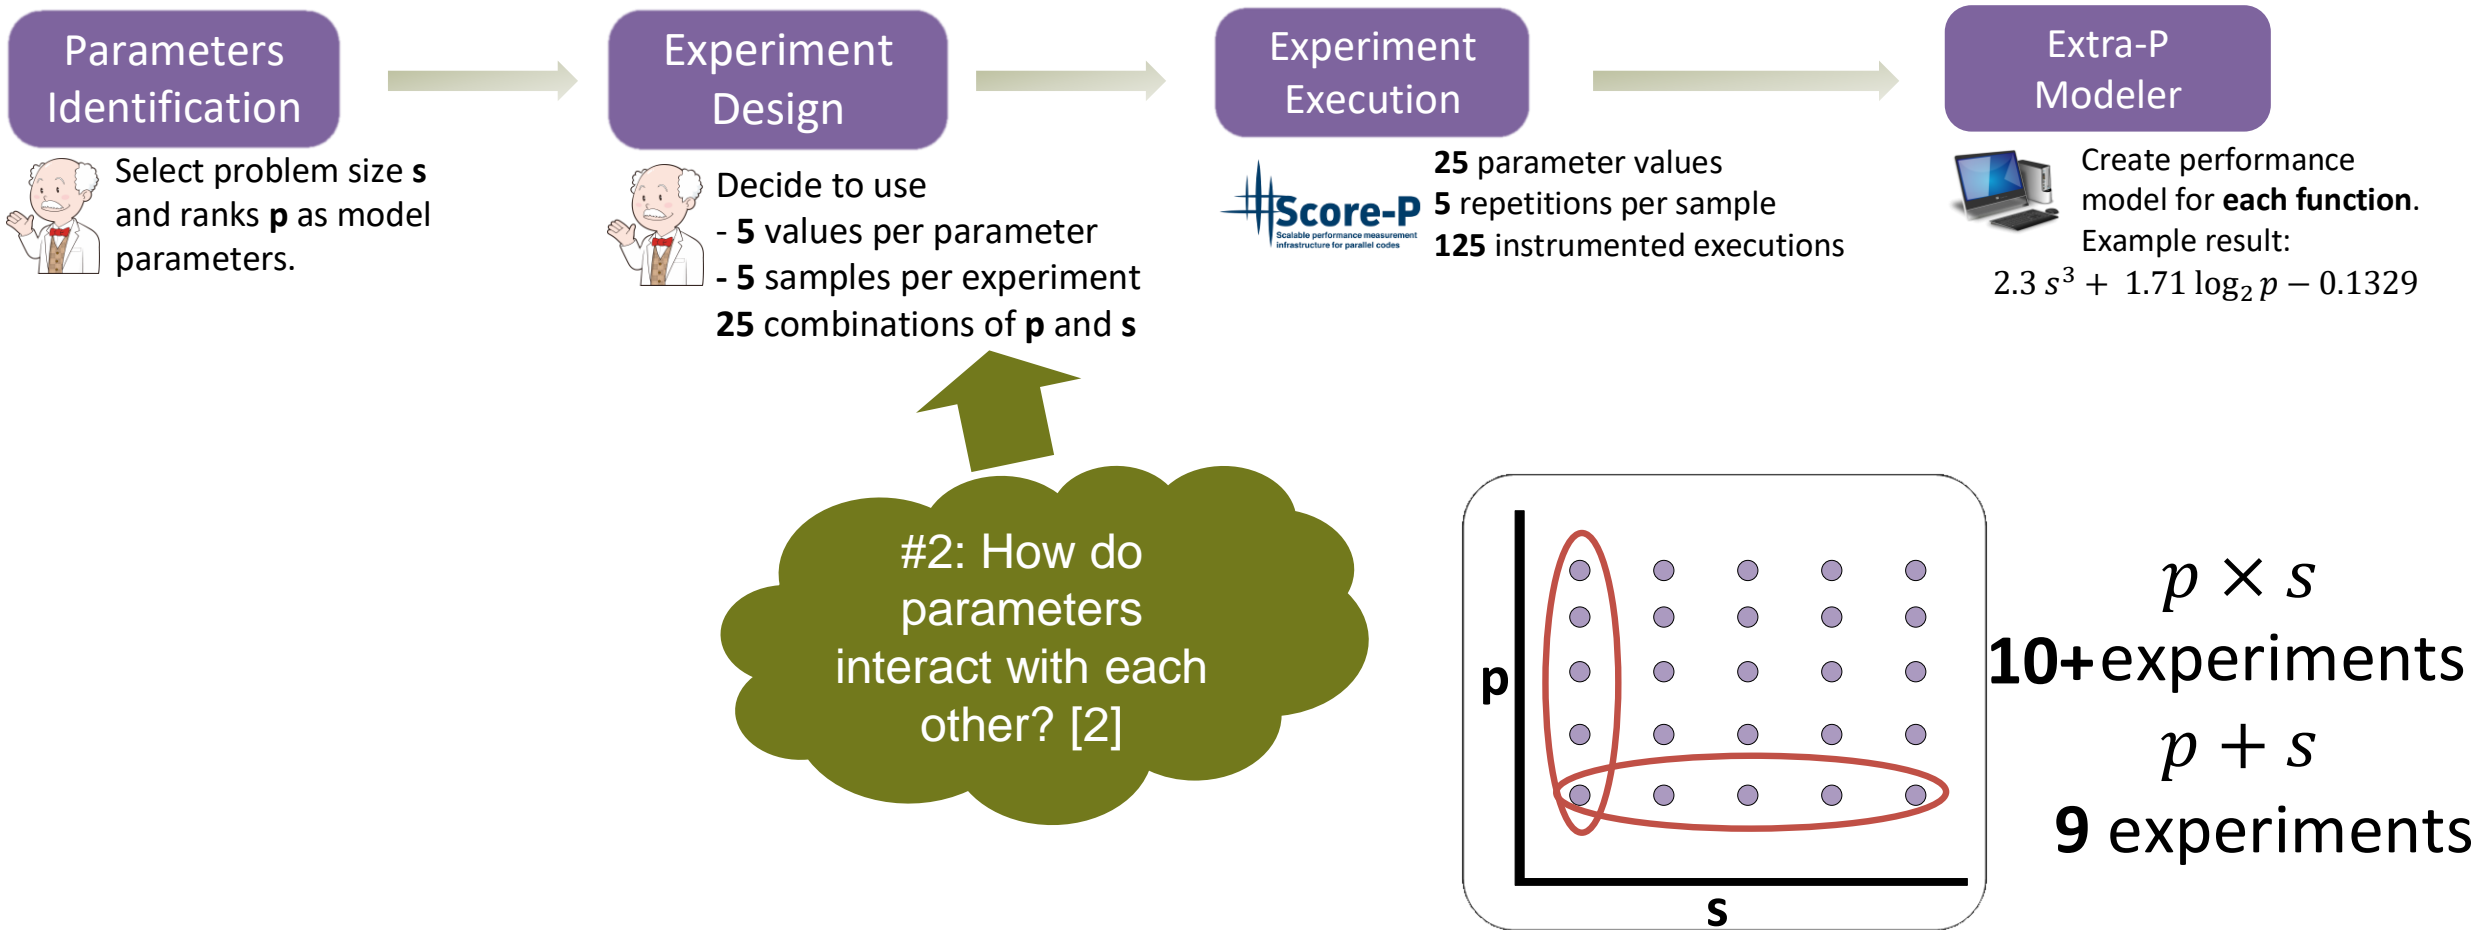
#1: Which parameters are relevant? [1]

```


int nx, ny, nz, nt;
int node_geometry[4];
int nflavors, propinterval;
int warms, trajecs, steps;
int niter, nrestart, prec_pbp;
  
```

A **subset** of all *su3_rmd* parameters.

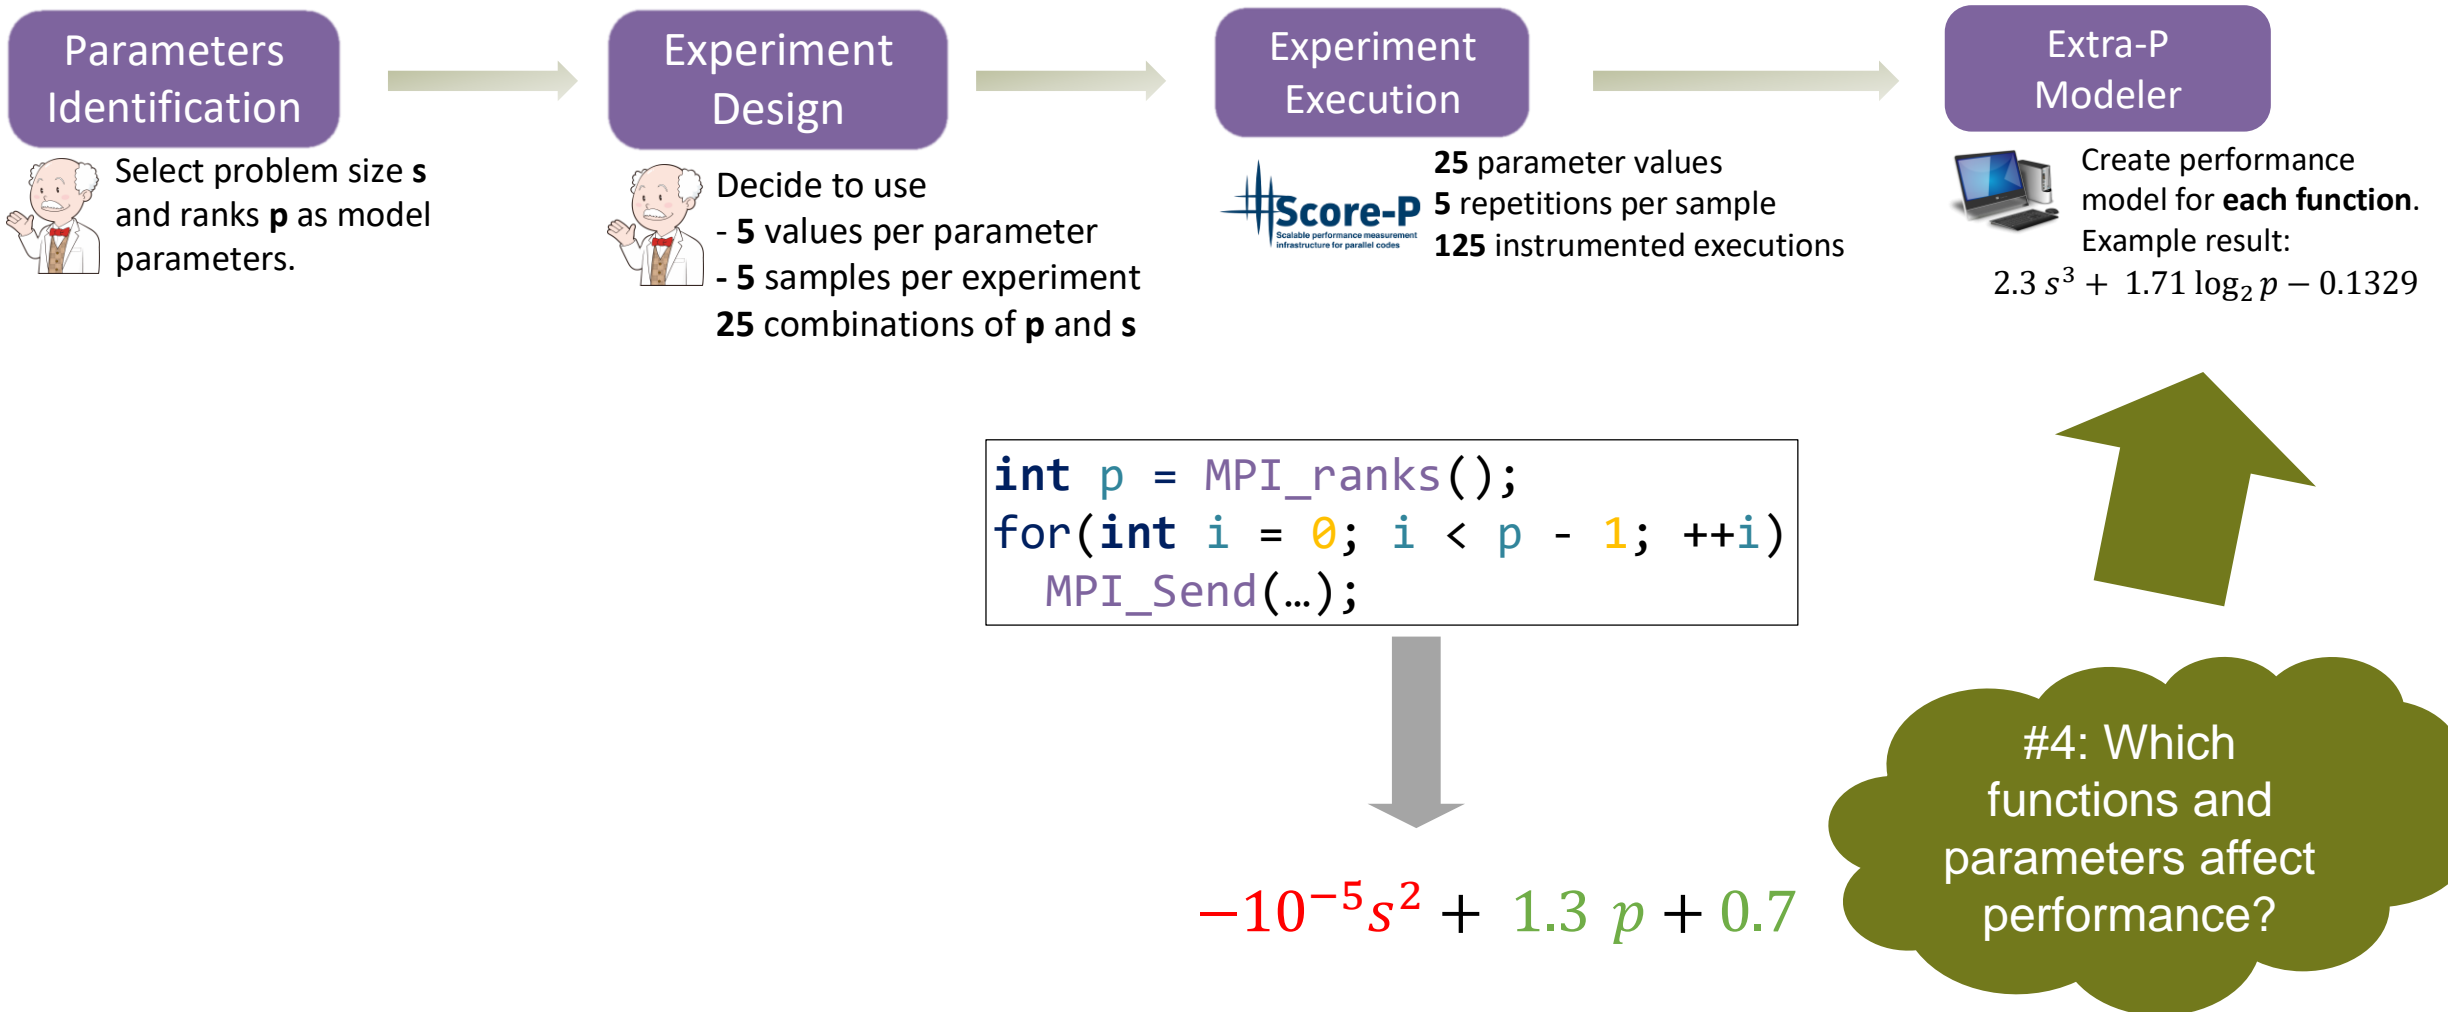
Challenges in Automatic Performance Modeling



Challenges in Automatic Performance Modeling



Challenges in Automatic Performance Modeling



Challenges in Automatic Performance Modeling

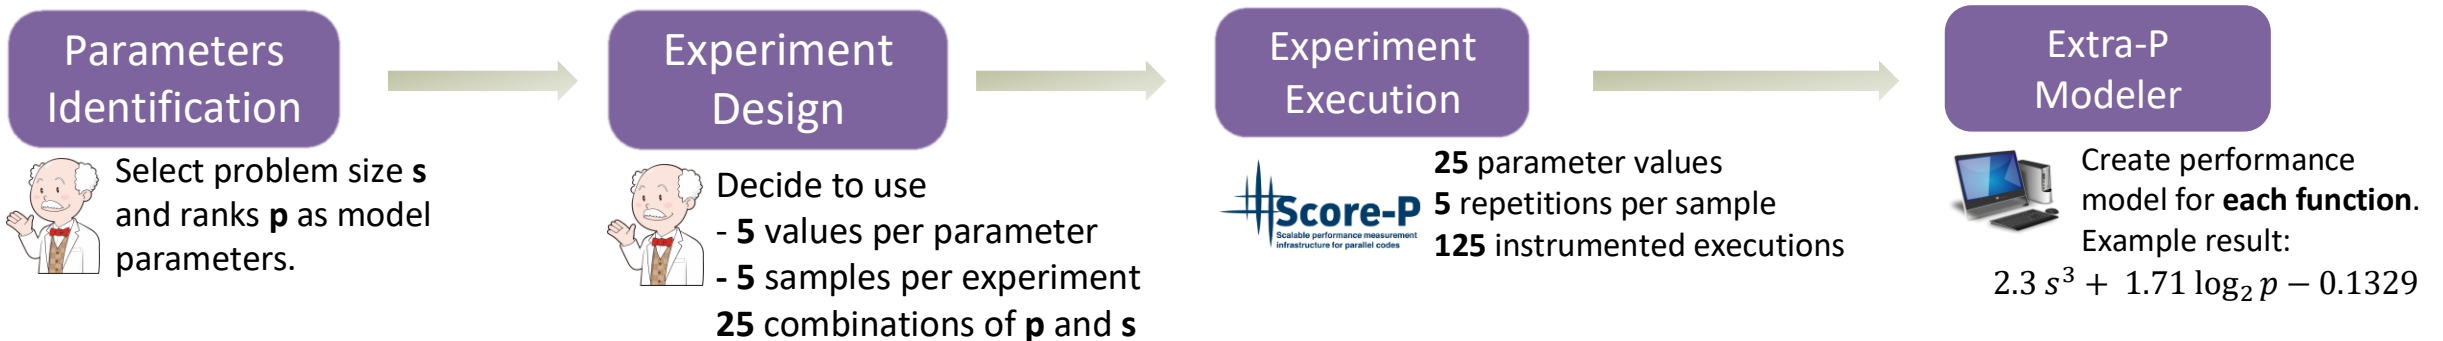


```
int p = MPI_ranks();
for(int i = 0; i < p - 1; ++i)
  MPI_Send(...);
```

#5: Does the model represent application behavior or hardware effects?

$$-10^{-5} s^2 + 1.3 p + 0.7$$

Challenges in Automatic Performance Modeling



#1: Which parameters are relevant? [1]

#2: How do parameters interact with each other? [2]

#3: Which functions are worth being instrumented?

#4: Which functions and parameters affect performance?

#5: Does the model represent application behavior or hardware effects?

Challenges in Automatic Performance Modeling



We need a **white-box approach**.

#1: Which parameters are relevant? [1]

other? [2]

#3: Which

#4: Which functions and parameters affect performance?

#5: Does the model represent application behavior or hardware effects?

What is important in our program?

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void i(int s, int p) {
    printf("%d %d\n", s, p);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

What is important in our program?

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

Which functions are performance-critical?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```


What is important in our program?

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

Which functions are performance-critical?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

What is important in our program?

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

Which functions are performance-critical?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

What is important in our program?

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

Which functions are performance-critical?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

What is important in our program?

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

Which functions are performance-critical?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

What is important in our program?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

Which functions are performance-critical?

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

What is important in our program?

```
void g(int s) {
    MPI_Send(&s, 1, MPI_INT);
}
```

```
void main(int s, int p) {
    g(s, p); h(s, p); i(s, p);
}
```

```
void h(int s, int p) {
    j(s);
}
```

```
void j(int x) {
    for(int j = 0; j < x; ++j)
        // compute
}
```

Which functions are performance-critical?

```
void i(int s, int p) {
    printf("%d\n", s, p);
}
```

Which parameters affect performance?

Analysis Requirements

The analysis must...

Analysis Requirements

The analysis must...

- ...detect functions which performance does not change.

Analysis Requirements

The analysis must...

- ...detect functions which performance does not change.
- ...detect which parameters affected non-constant loops.

Analysis Requirements

The analysis must...

- ...detect functions which performance does not change.
- ...detect which parameters affected non-constant loops.
- ...support inter-procedural parameter dependencies.

Analysis Requirements

The analysis must...

- ...detect functions which performance does not change.
- ...detect which parameters affected non-constant loops.
- ...support inter-procedural parameter dependencies.
- ...be memory agnostic.

Analysis Requirements

The analysis must...

- ...detect functions which performance does not change.
- ...detect which parameters affected non-constant loops.
- ...support inter-procedural parameter dependencies.
- ...be memory agnostic.
- ...support parallelism patterns introduced by MPI.

Analysis Requirements

The analysis must...

- ...detect functions which performance does not change.
- ...detect which parameters affected non-constant loops.
- ...support inter-procedural parameter dependencies.
- ...be memory agnostic.
- ...support parallelism patterns introduced by MPI.

Answer? Taint Analysis.

Taint Analysis: track parameters propagation

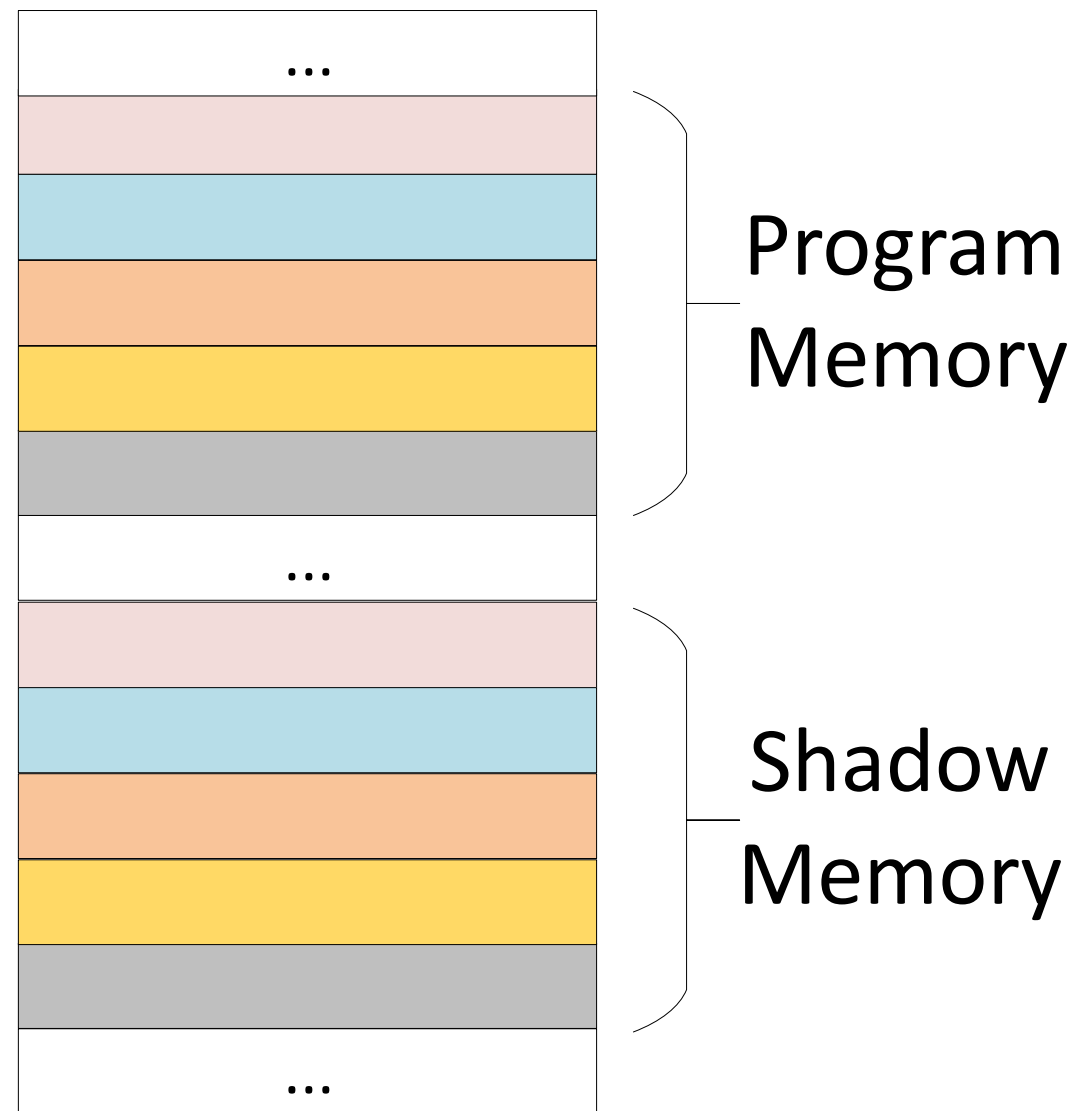
```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```



Taint Analysis: track parameters propagation

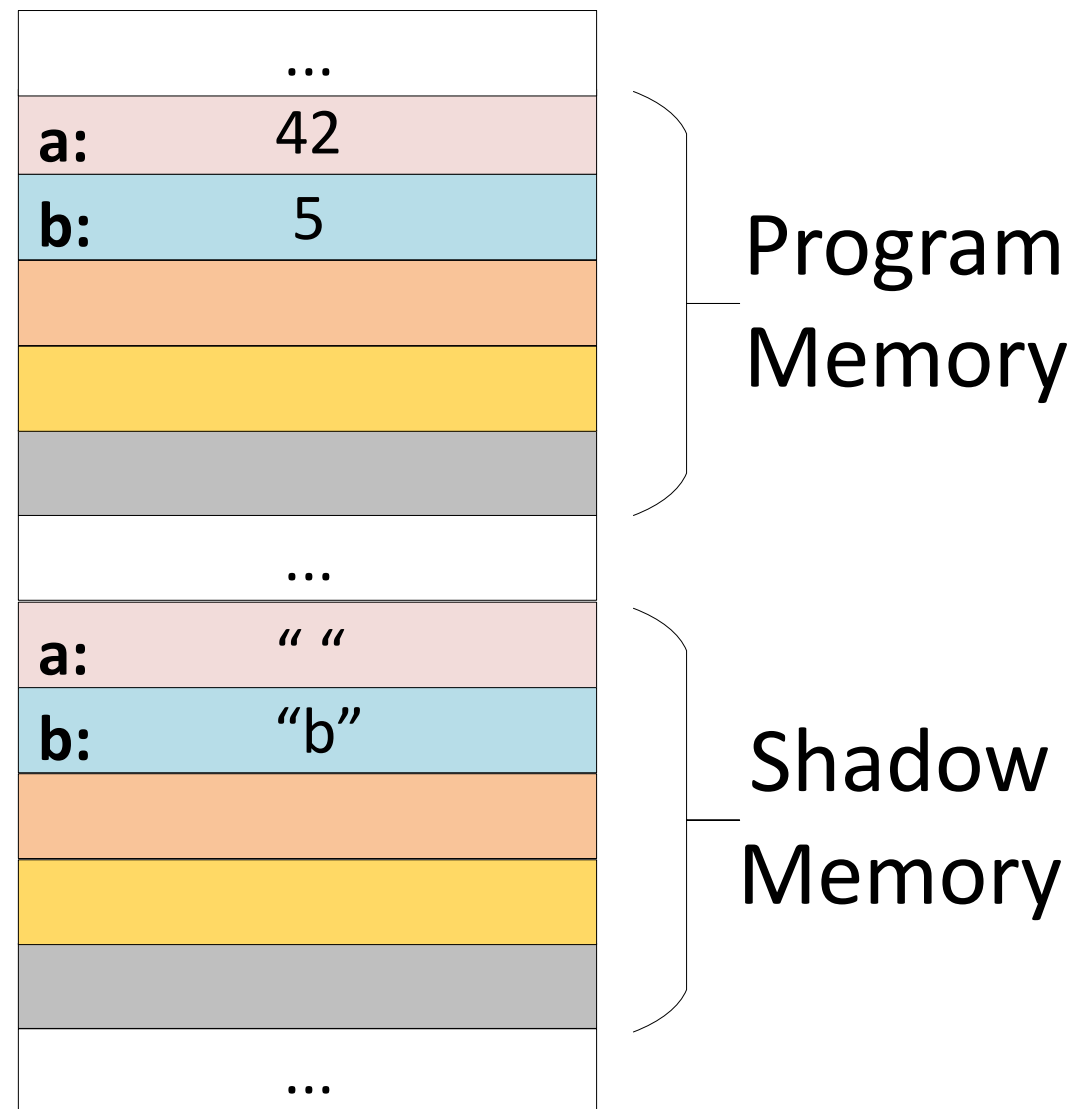
```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```



Taint Analysis: track parameters propagation

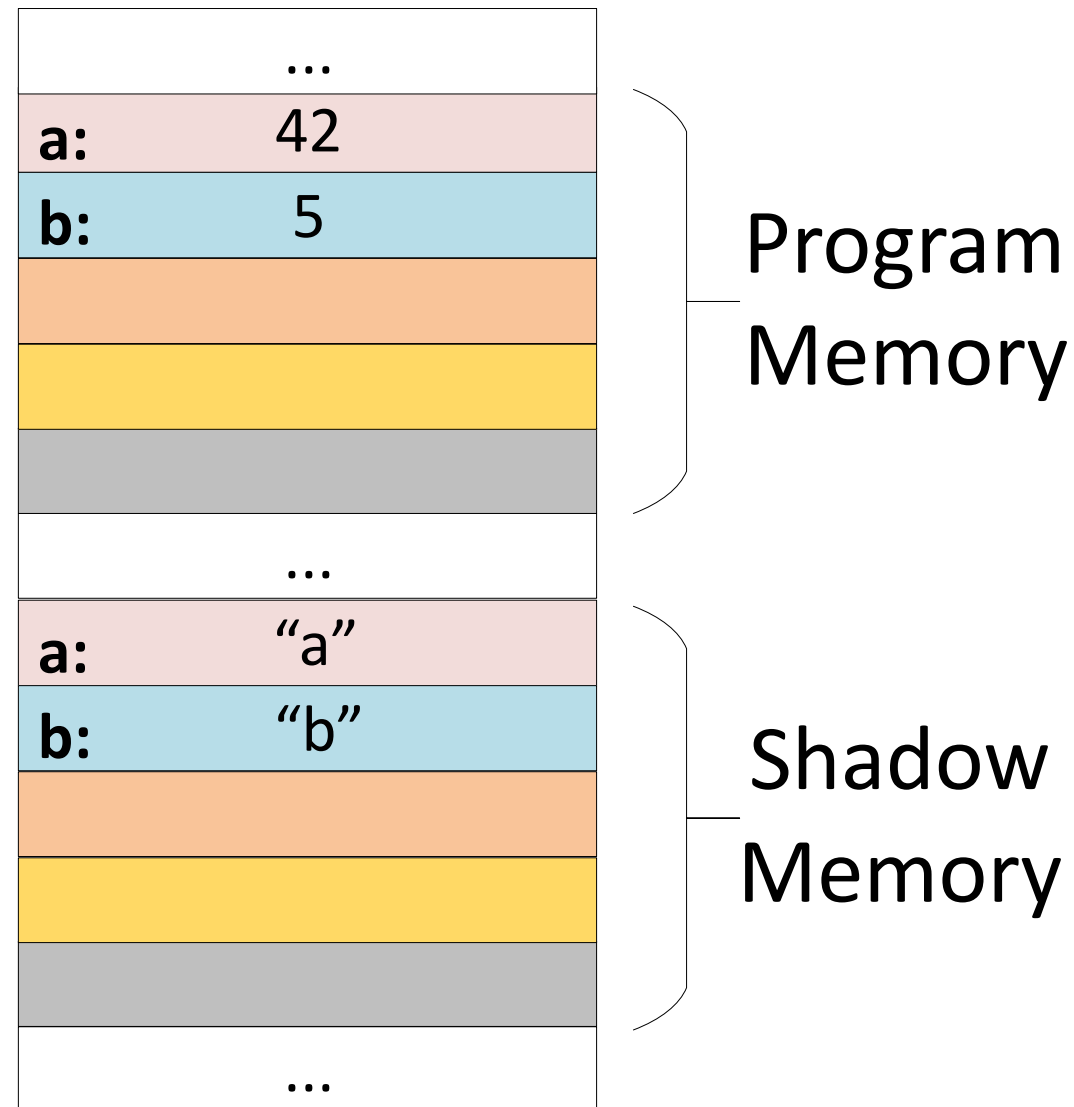
```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```



Taint Analysis: track parameters propagation

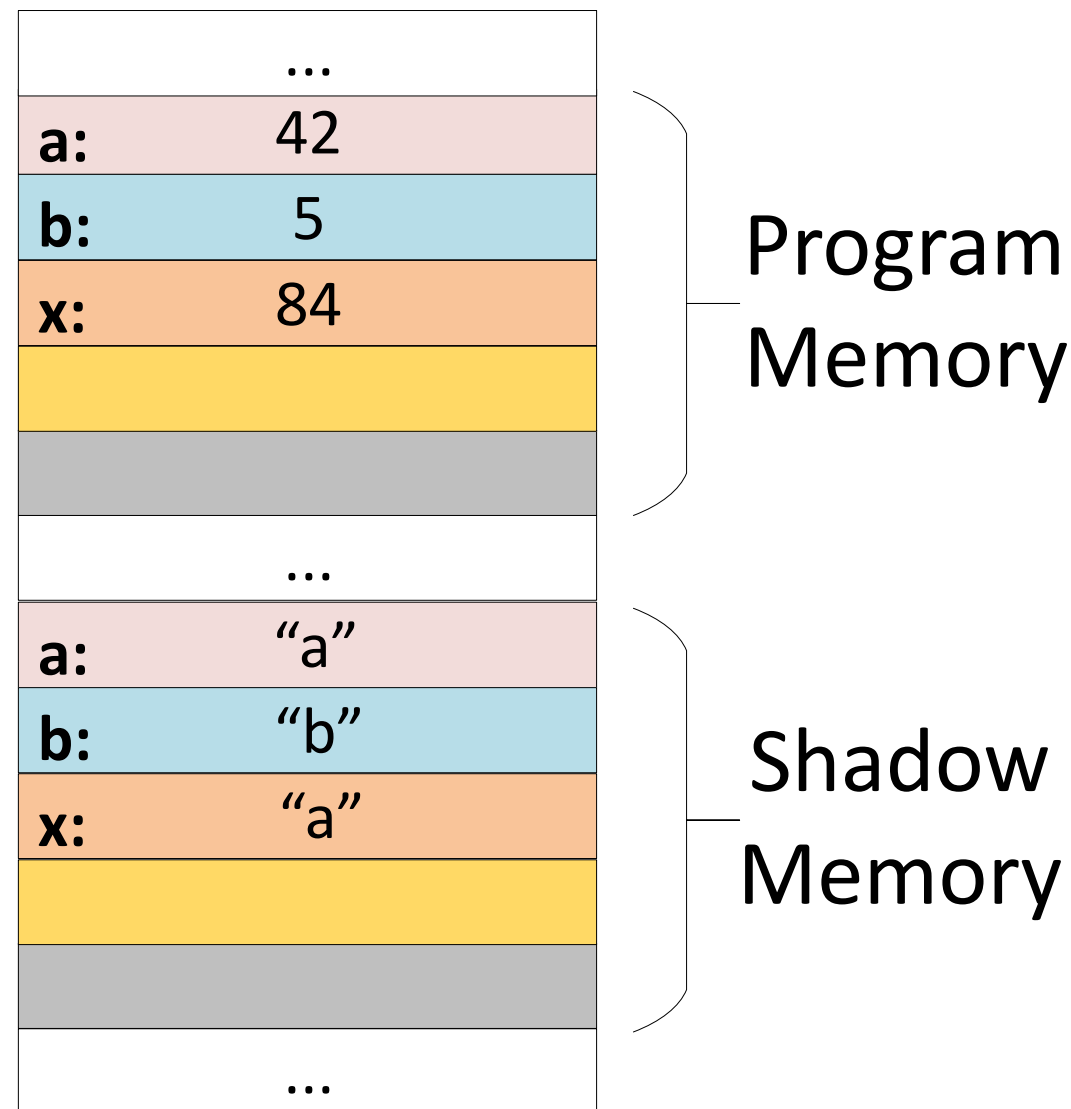
```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```



Taint Analysis: track parameters propagation

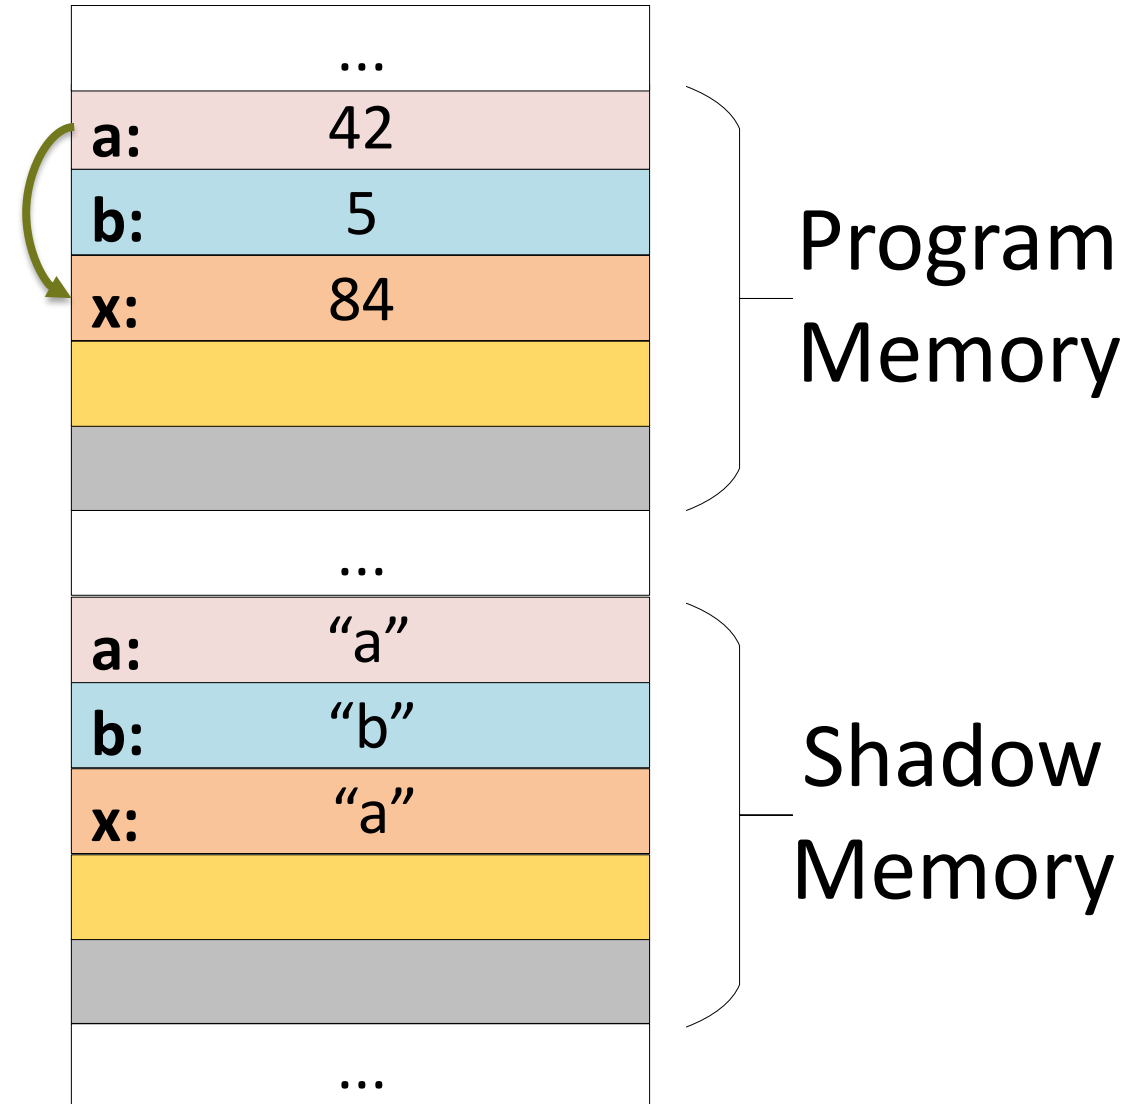
```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```



Taint Analysis: track parameters propagation

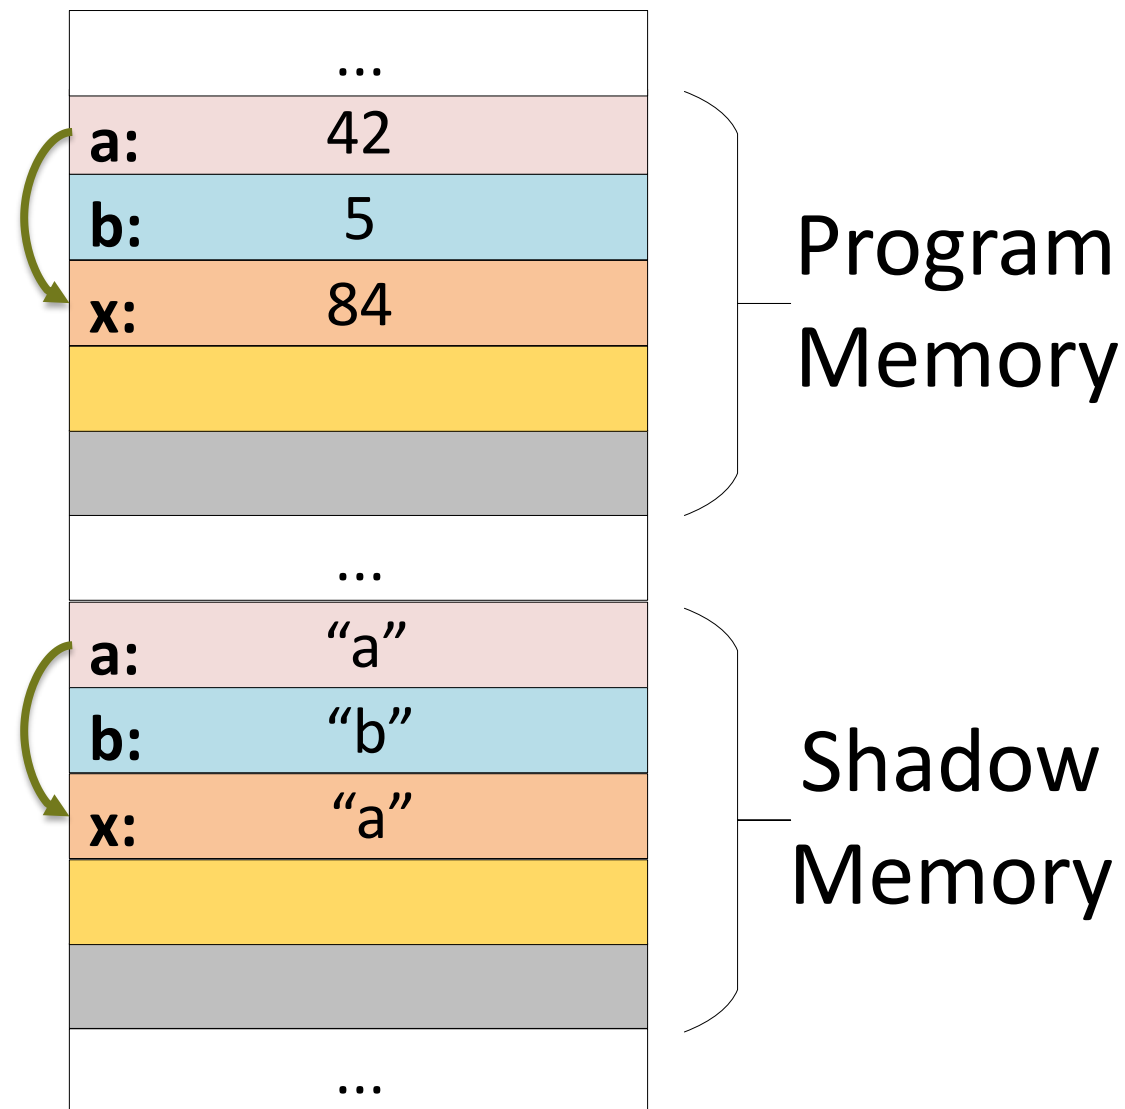
```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```



Taint Analysis: track parameters propagation

```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
```

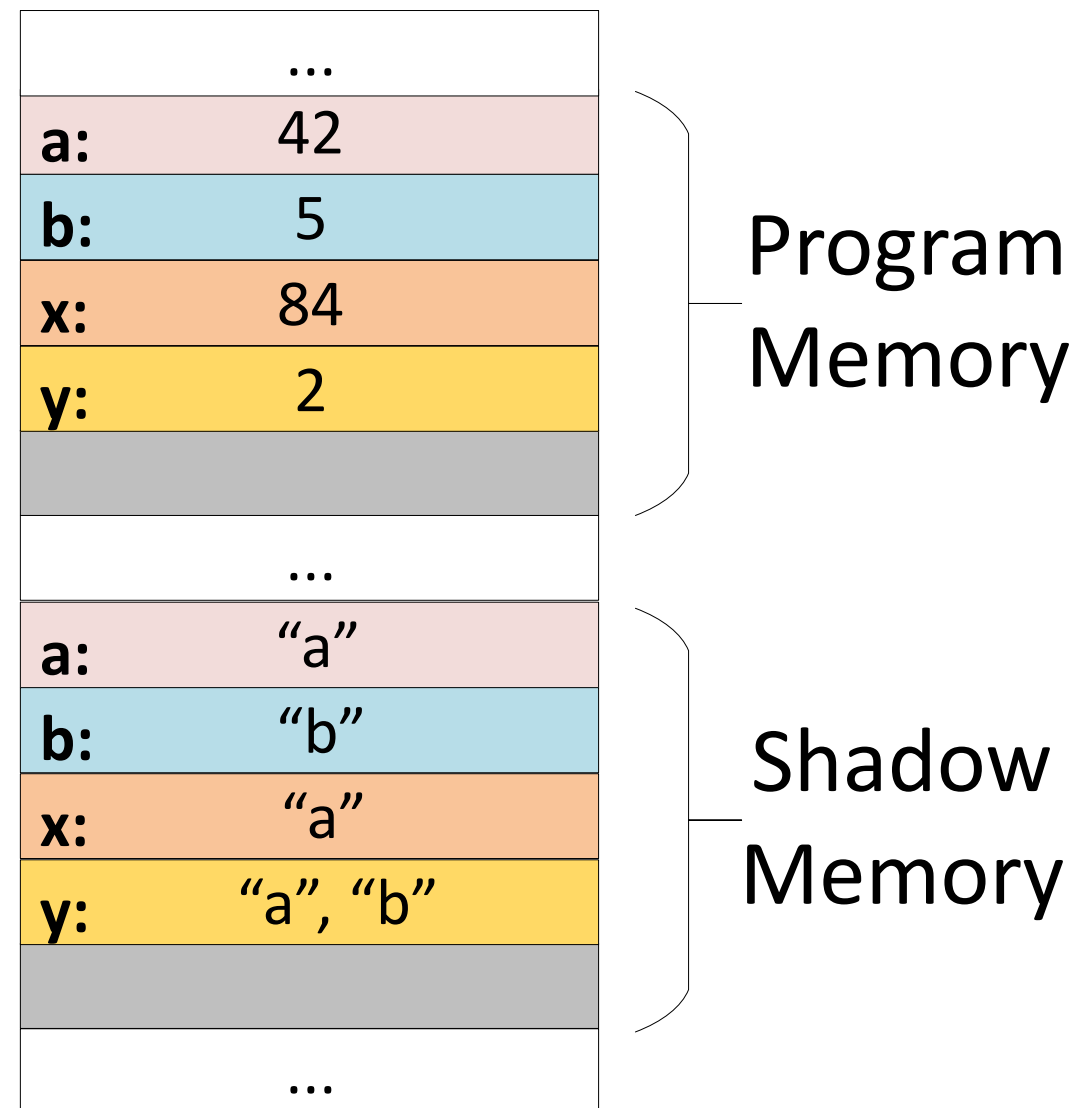
```
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
```

```
if(a != 43)
```

```
    z = 6;
```



Taint Analysis: track parameters propagation

```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

// Control-flow propagation

```
int z = 10;
if(a != 43)
    z = 6;
```

...		Program Memory
a:	42	
b:	5	
x:	84	
y:	2	
z:	10	Shadow Memory
...		
a:	"a"	
b:	"b"	
x:	"a"	
y:	"a", "b"	
...		

Taint Analysis: track parameters propagation

```
int a = 42;
int b; MPI_Comm_size(&b, comm);
taint_variable(a);
```

// Data-flow propagation

```
int x = 2 * a;
int y = modulo(a, b);
```

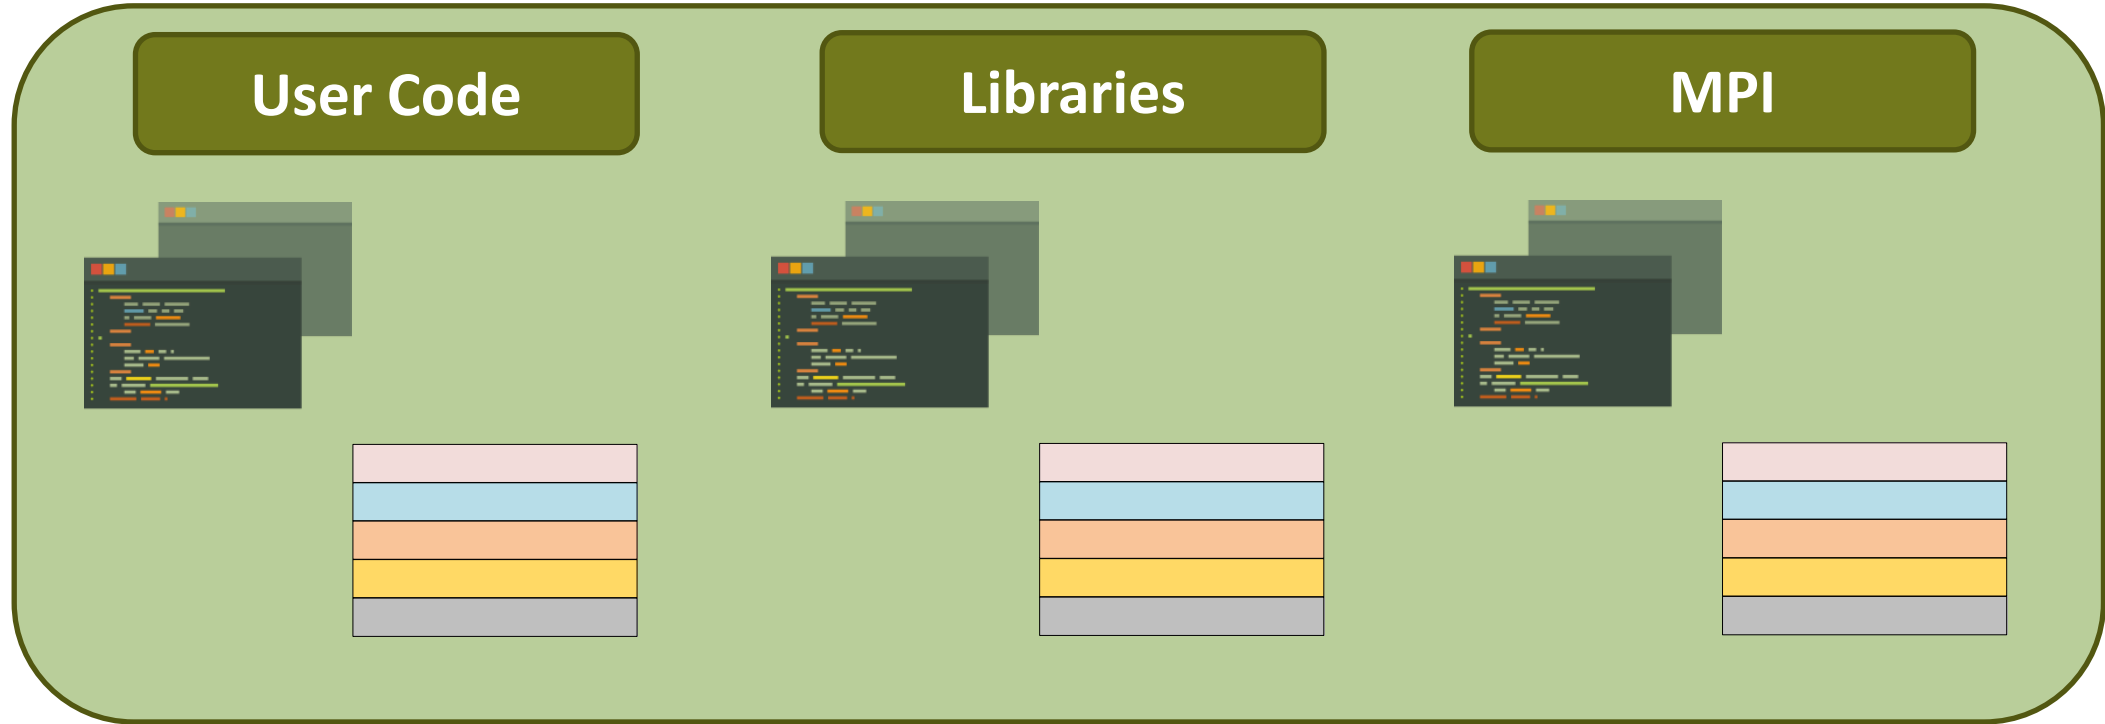
// Control-flow propagation

```
int z = 10;
if(a != 43)
```

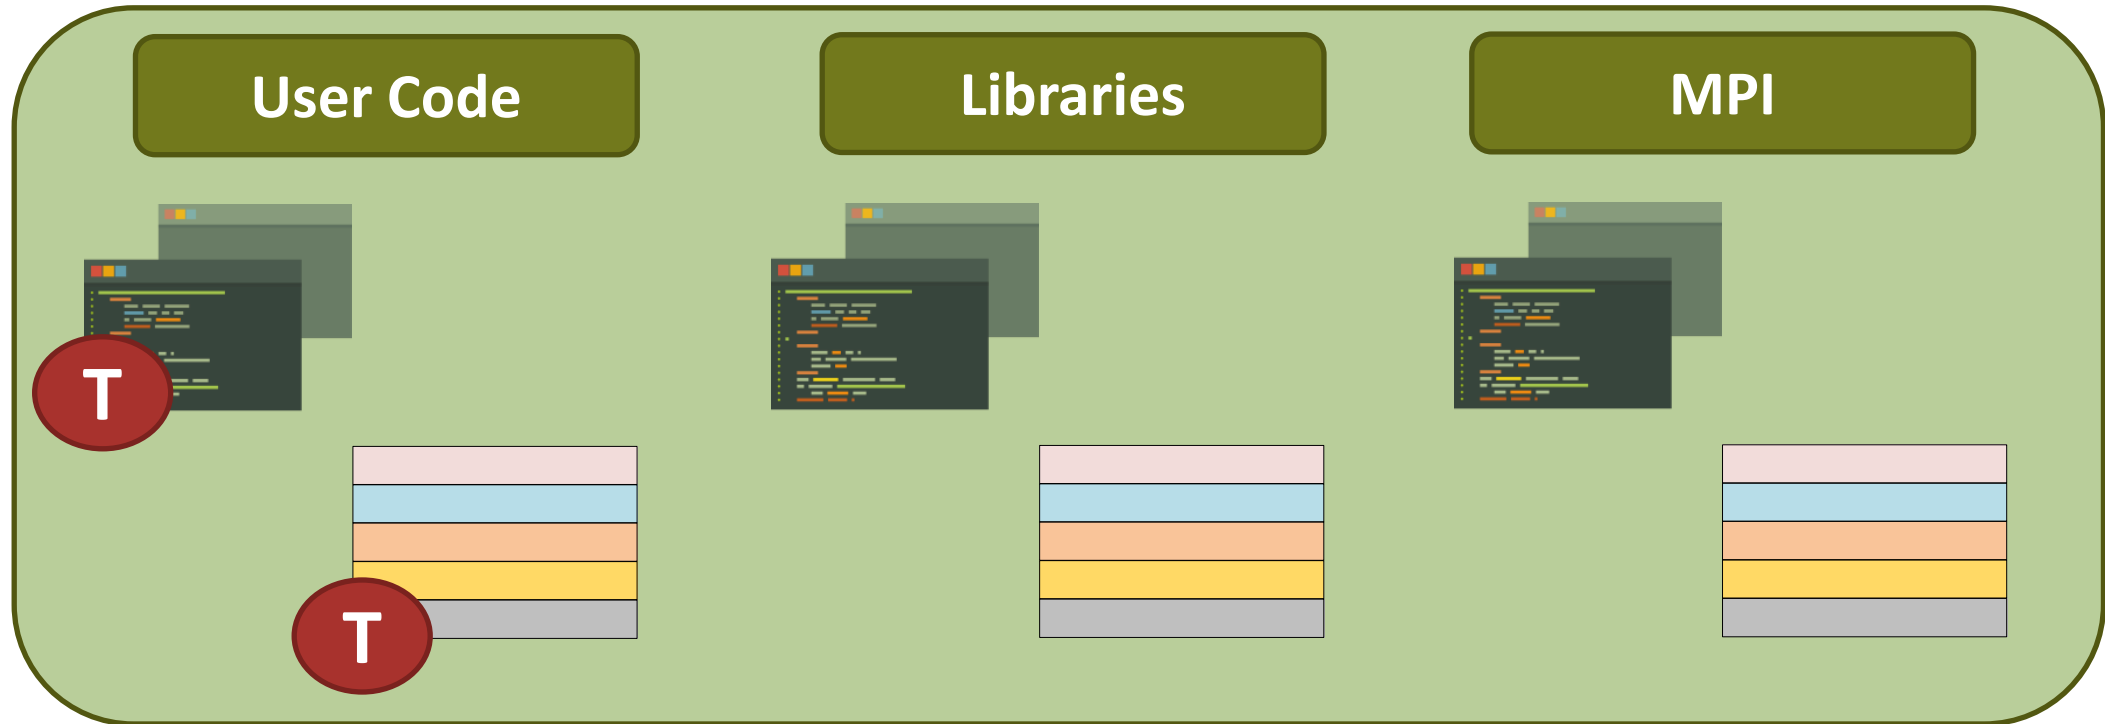
```
    z = 6;
```

...		Program Memory
a:	42	
b:	5	
x:	84	
y:	2	
z:	6	
...		Shadow Memory
a:	"a"	
b:	"b"	
x:	"a"	
y:	"a", "b"	
z:	"a"	
...		


Support for MPI



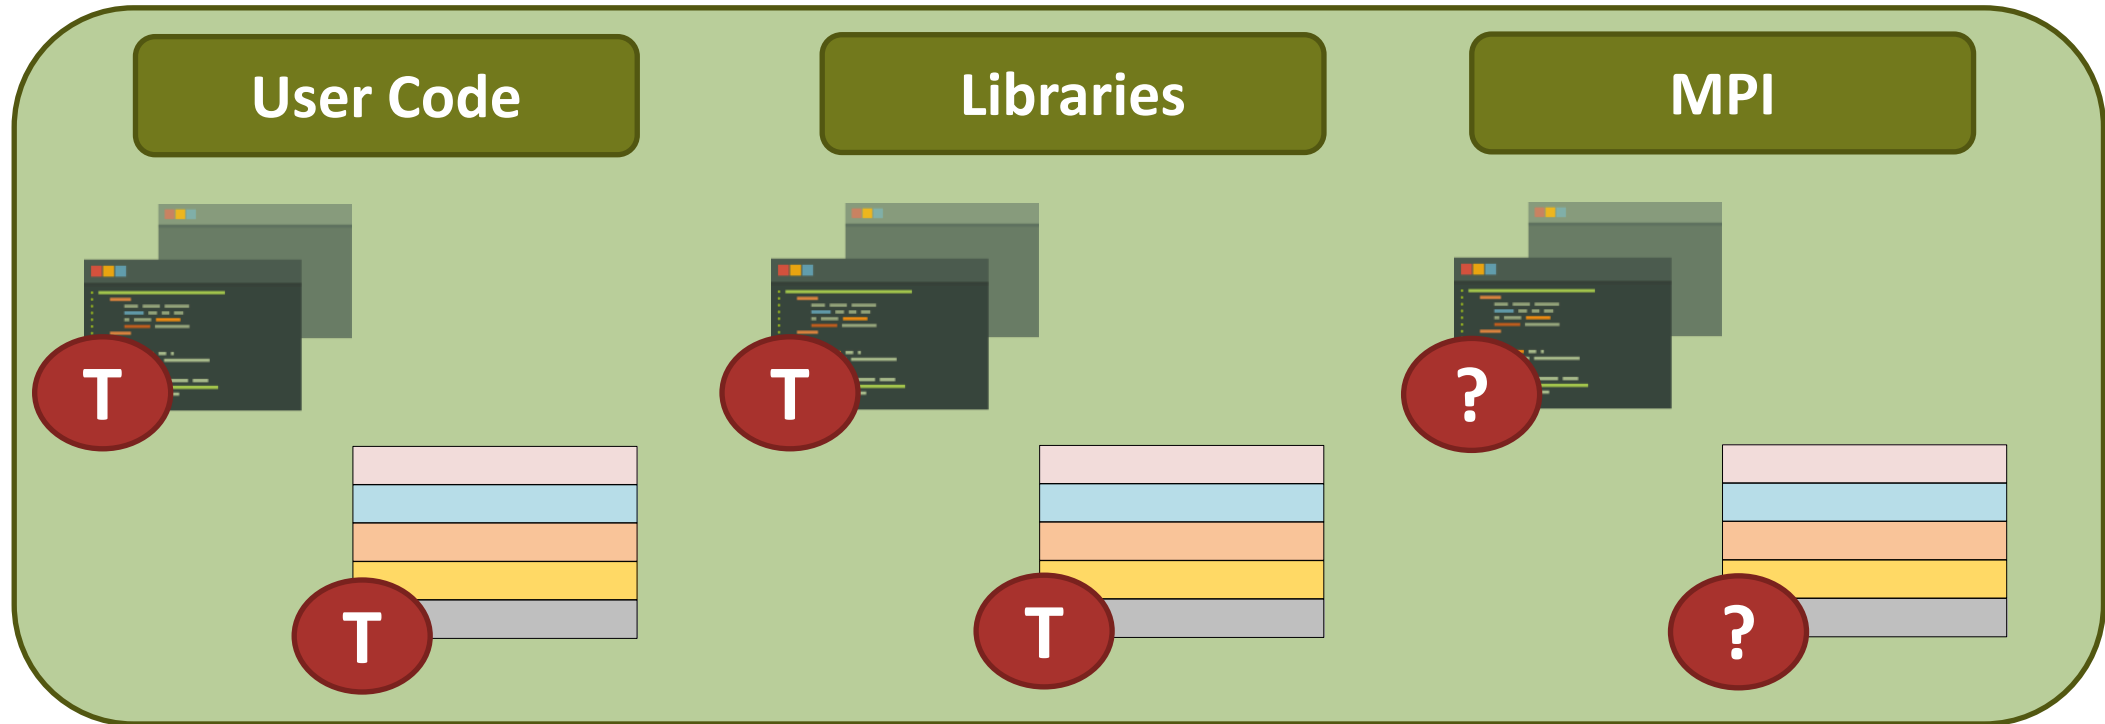
Support for MPI




Support for MPI



Support for MPI



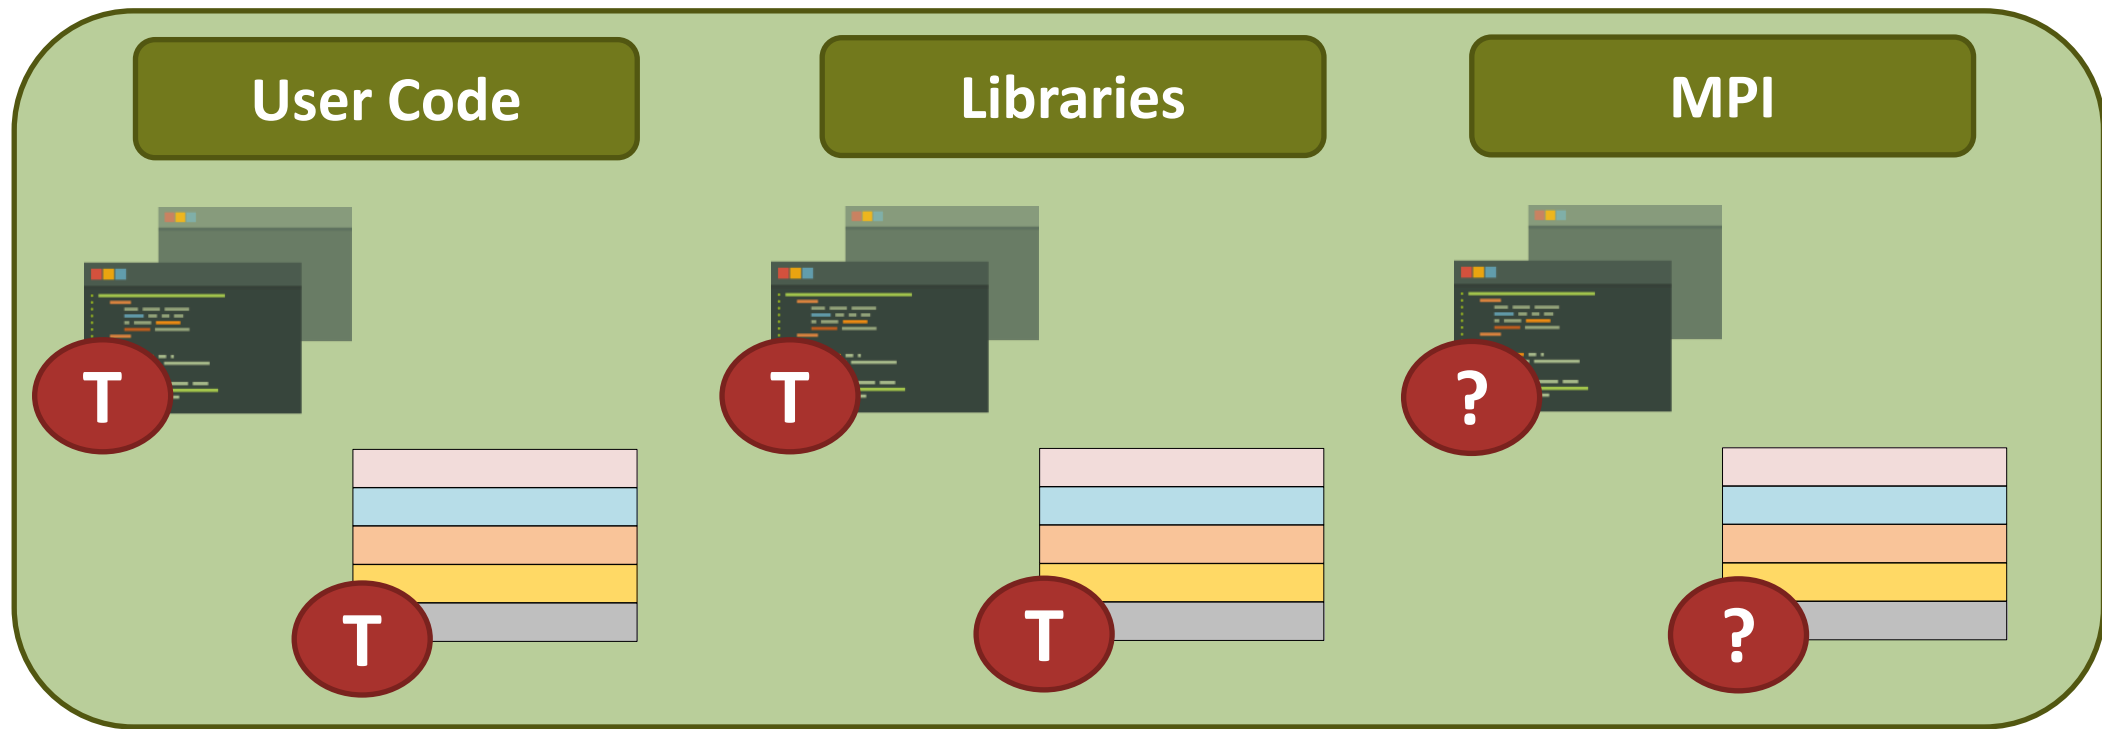
Support for MPI



Compile MPI with taint support?

- ✗ No support for closed implementations.
- ✗ Troublesome deployment on a cluster.

Support for MPI



Compile MPI with taint support?

Define taint characteristics of MPI.

- ✗ No support for closed implementations.
- ✗ Troublesome deployment on a cluster.

- ✓ No recompilation needed.
- ✓ Works with every implementation.

Support for MPI

Taint Labels

`p` - # of MPI ranks

Taint Sources

`MPI_Comm_size(MPI_Comm, int *)`

Taint Sinks

`MPI_Reduce(const void *, void *, int count, MPI_Datatype, MPI_Op, int root, MPI_Comm);`

Support for MPI

Taint Labels

`p` - # of MPI ranks

Taint Sources

`MPI_Comm_size(MPI_Comm, int *)`

Taint Sinks

`MPI_Reduce(const void *, void *, int count, MPI_Datatype, MPI_Op, int root, MPI_Comm);`

Support for MPI

Taint Labels

`p` - # of MPI ranks

Taint Sources

`MPI_Comm_size(MPI_Comm, int *)`

Taint Sinks

`MPI_Reduce(const void *, void *, int count, MPI_Datatype, MPI_Op, int root, MPI_Comm);`

Support for MPI

Taint Labels

`p` - # of MPI ranks

Taint Sources

`MPI_Comm_size(MPI_Comm, int *)`


Taint Sinks

`MPI_Reduce(const void *, void *, int count, MPI_Datatype, MPI_Op, int root, MPI_Comm);`

perf-taint: Hybrid Taint Analysis



perf-taint: Hybrid Taint Analysis



Annotate Parameters


```
register_variable("size", &size);
```

perf-taint: Hybrid Taint Analysis




Annotate Parameters

`register_variable("size", &size);`




Static Loop
Analysis

perf-taint: Hybrid Taint Analysis



Annotate Parameters

`register_variable("size", &size);`




Static Loop
Analysis



Dynamic Taint
Analysis

perf-taint: Hybrid Taint Analysis



Annotate Parameters

`register_variable("size", &size);`



Static Loop
Analysis



Dynamic Taint
Analysis



(1) Parametric Dependencies
(2) Constant Functions

How do we apply this knowledge?

How do we apply this knowledge?

Parameters Identification



Select problem size s
and ranks p as model
parameters.

How do we apply this knowledge?

Parameters Identification



Select problem size s
and ranks p as model
parameters.



**Expert selects
parameters.**

How do we apply this knowledge?

Parameters Identification



Select problem size s and ranks p as model parameters.



Expert selects parameters.



Taint-based coverage selects parameters.

How do we apply this knowledge?

Parameters Identification



Experiment Design



Select problem size s and ranks p as model parameters.



Decide to use

- 5 values per parameter
- 5 samples per experiment

25 combinations of p and s



Expert selects parameters.



Taint-based coverage selects parameters.

How do we apply this knowledge?

Parameters Identification



Experiment Design



Select problem size s and ranks p as model parameters.



Decide to use
- 5 values per parameter
- 5 samples per experiment
25 combinations of p and s



Expert selects parameters.



Use complex heuristics.



Taint-based coverage selects parameters.

How do we apply this knowledge?

Parameters Identification



Experiment Design



Select problem size s and ranks p as model parameters.



Decide to use

- 5 values per parameter
- 5 samples per experiment

25 combinations of p and s



Expert selects parameters.



Use complex heuristics.




Taint-based coverage selects parameters.






Use parameter dependencies.


How do we apply this knowledge?




 **Expert selects parameters.**


  **Use complex heuristics.**



 **Taint-based coverage selects parameters.**


 **Use parameter dependencies.**


How do we apply this knowledge?




 **Expert selects parameters.**


  **Use complex heuristics.**


 **Instrument all functions.**



 **Taint-based coverage selects parameters.**


 **Use parameter dependencies.**


How do we apply this knowledge?





 **Expert selects parameters.**

  **Use complex heuristics.**

 **Instrument all functions.**


 **Taint-based coverage selects parameters.**



 **Use parameter dependencies.**

 **Skip irrelevant functions.**


How do we apply this knowledge?




 Expert selects parameters.

  Use complex heuristics.

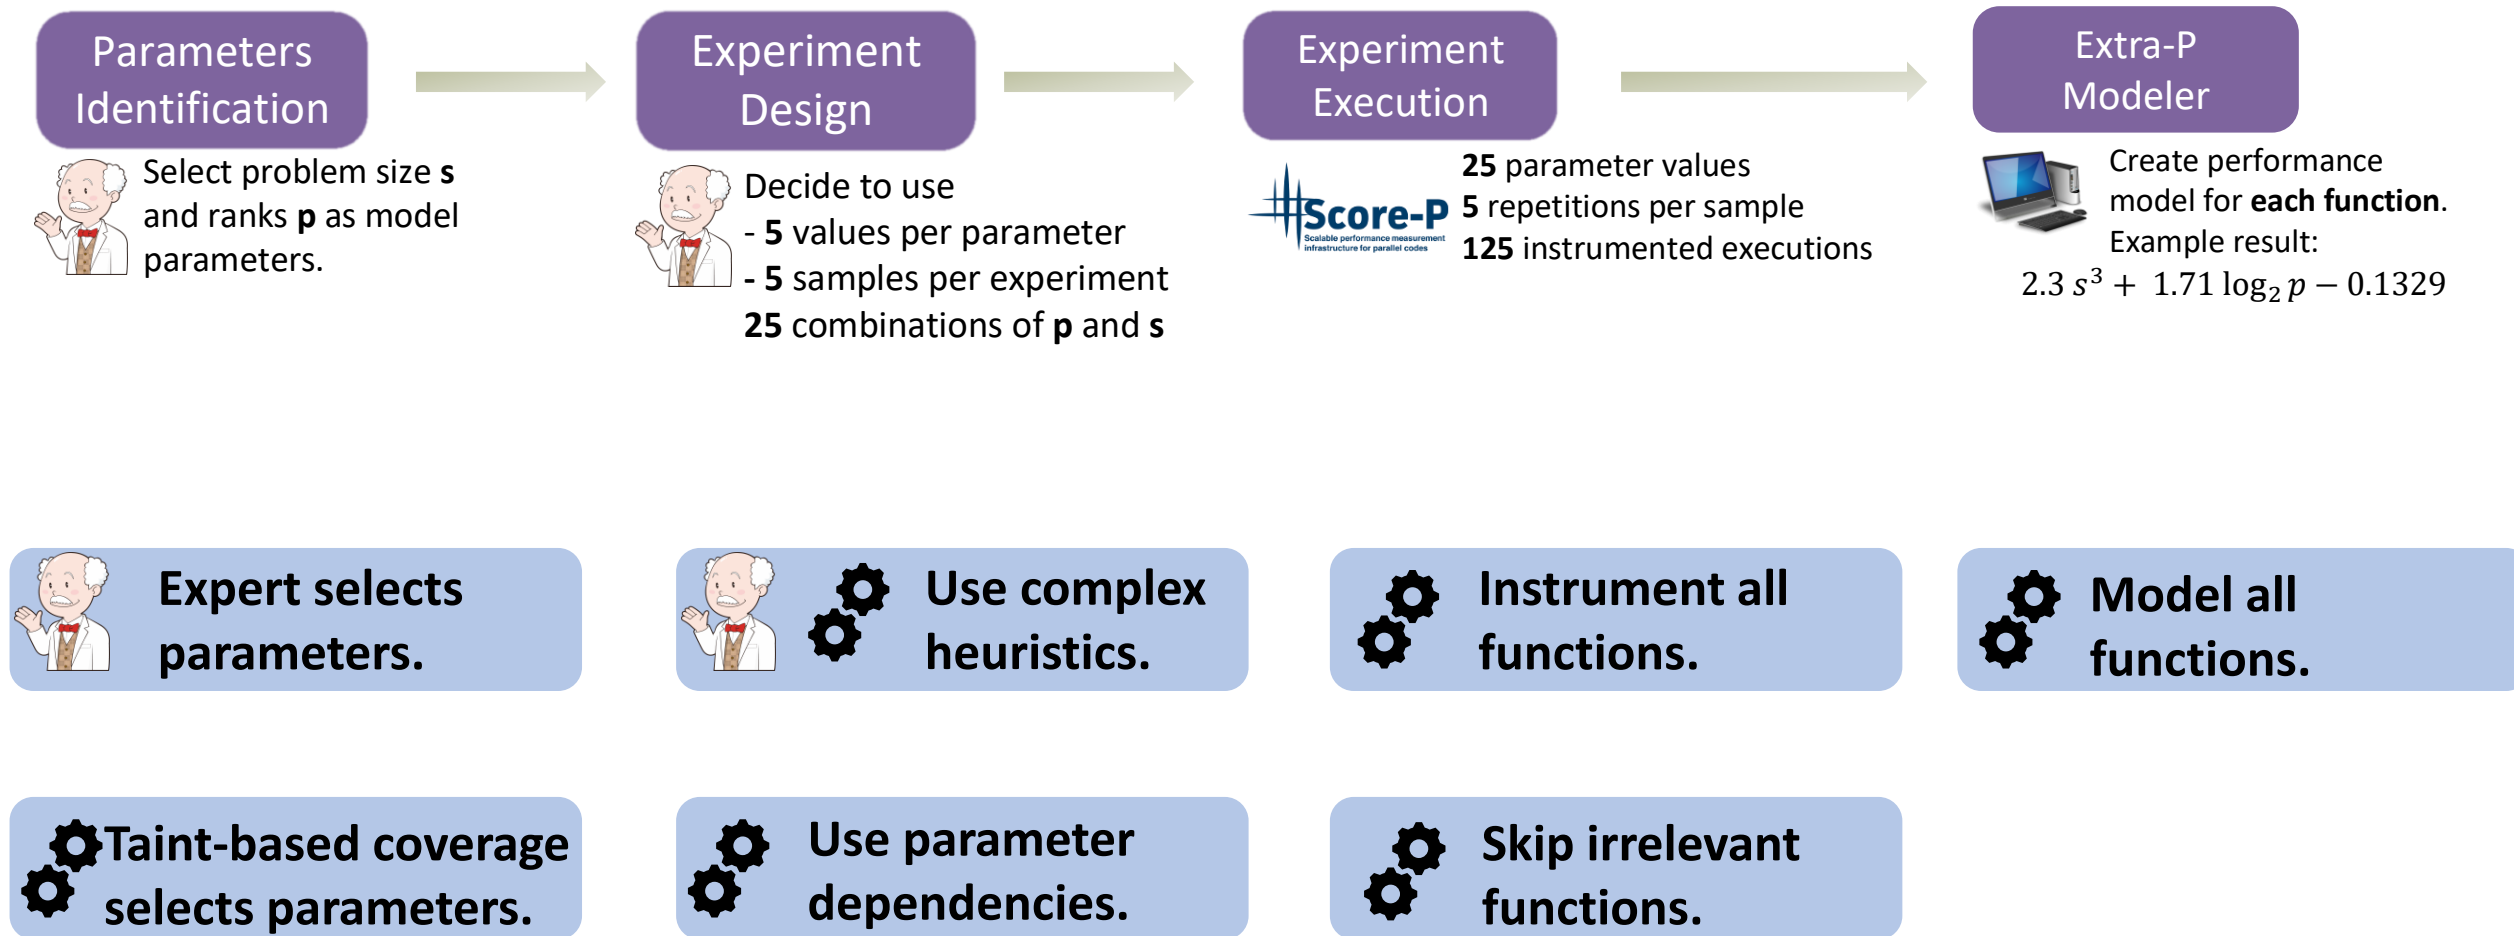
 Instrument all functions.

 Taint-based coverage selects parameters.

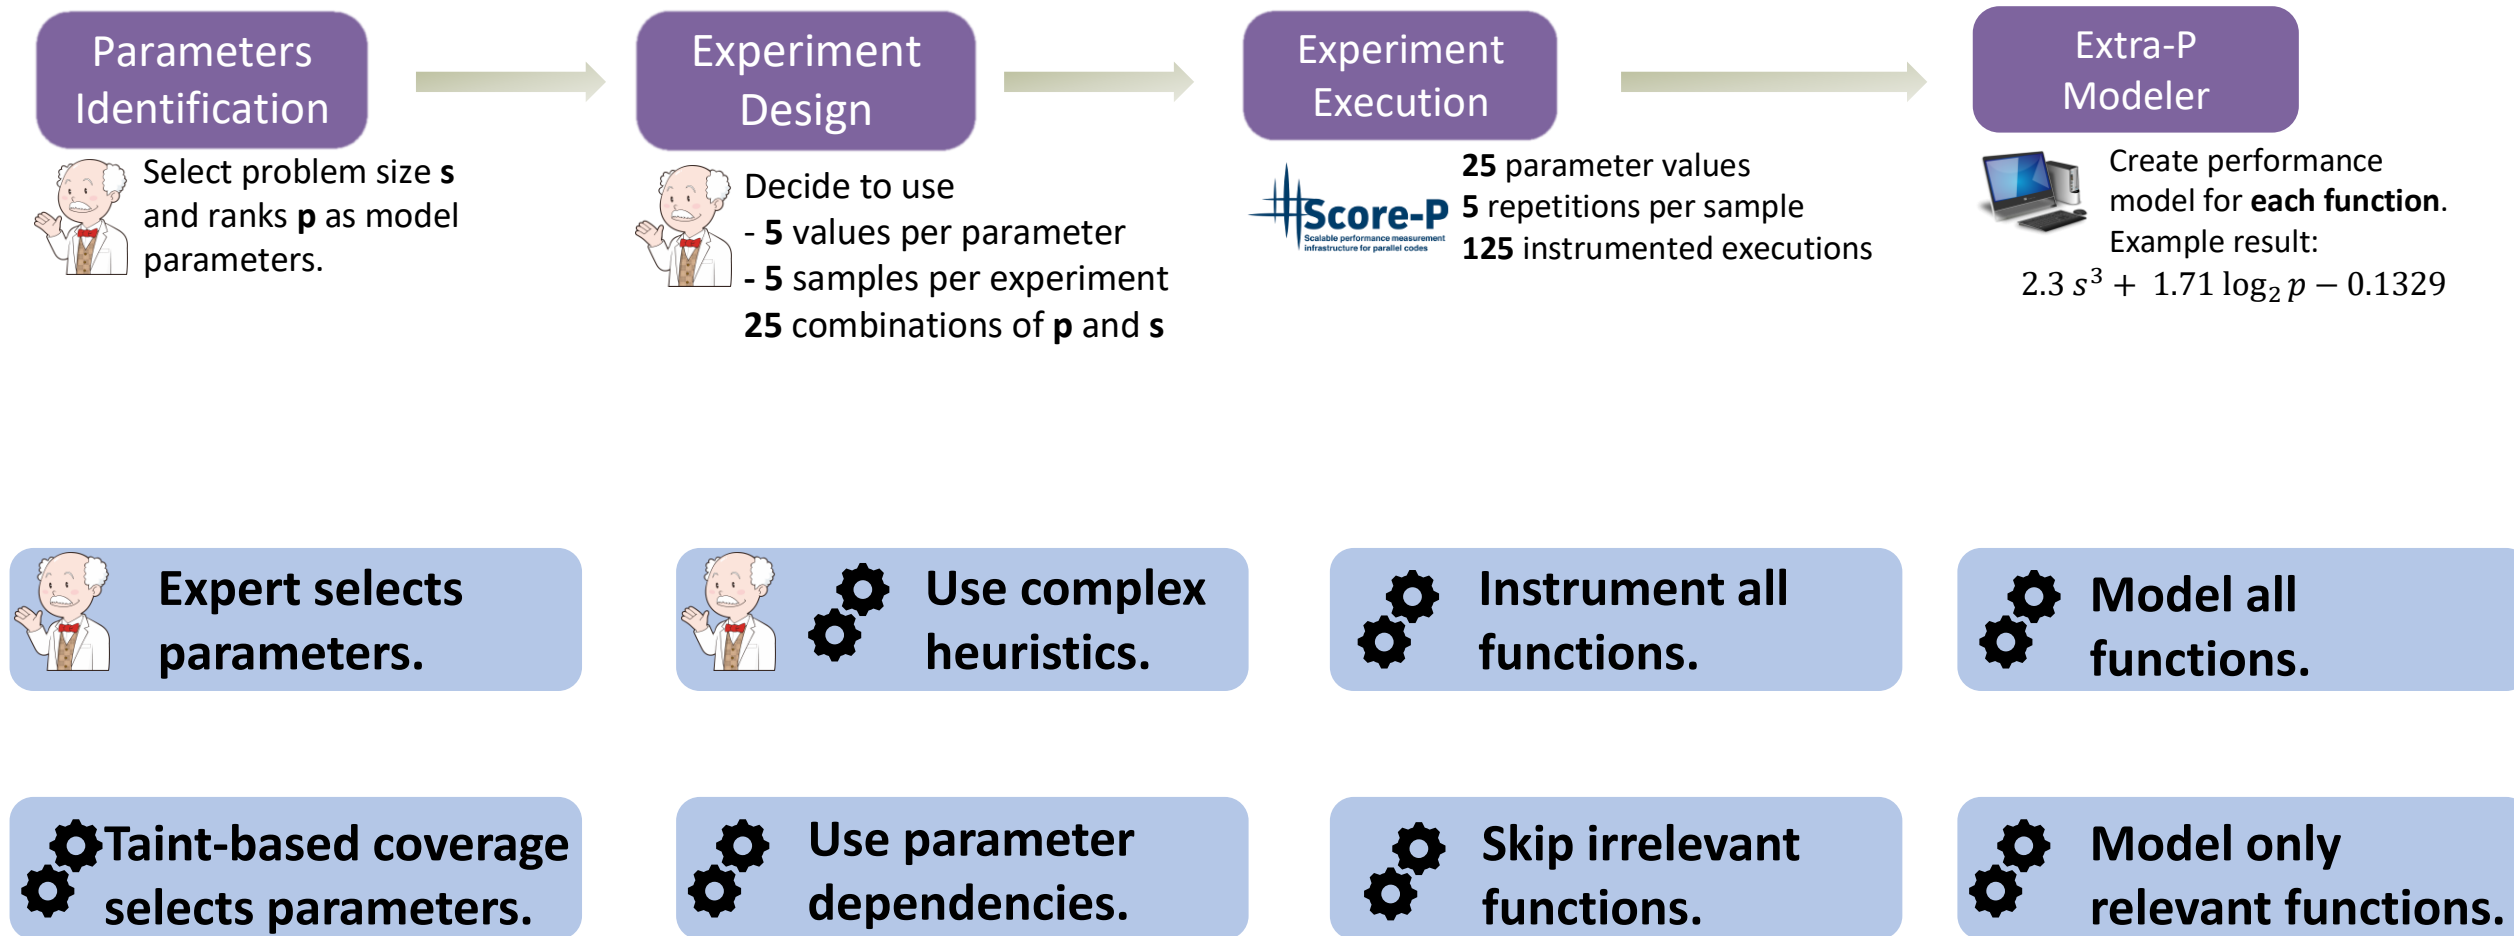
 Use parameter dependencies.

 Skip irrelevant functions.

How do we apply this knowledge?



How do we apply this knowledge?



Case Studies

Cost
Fewer Experiments
Cheaper Experiments

Case Studies

Cost

Fewer Experiments
Cheaper Experiments

Quality

Less Intrusion
More Noise Resilience

Case Studies

Cost

Fewer Experiments
Cheaper Experiments

Quality

Less Intrusion
More Noise Resilience

Validity

Experiment Design
Hardware Contention

Case Studies

Case Studies

Piz Daint, 21 nodes

- Intel Xeon E5-2695 v4 2.1 GHz
- 2 sockets, 18 cores each
- 128 GB Memory
- Cray MPICH 7.7.2

Case Studies

Piz Daint, 21 nodes

- Intel Xeon E5-2695 v4 2.1 GHz
- 2 sockets, 18 cores each
- 128 GB Memory
- Cray MPICH 7.7.2

Skylake Cluster, 2 nodes

- Intel Xeon 6154 3 GHz
- 36 cores
- 384 GB Memory
- OpenMPI 4.0.3

Case Studies

Piz Daint, 21 nodes

- Intel Xeon E5-2695 v4 2.1 GHz
- 2 sockets, 18 cores each
- 128 GB Memory
- Cray MPICH 7.7.2

Skylake Cluster, 2 nodes

- Intel Xeon 6154 3 GHz
- 36 cores
- 384 GB Memory
- OpenMPI 4.0.3

LULESH

- p : 27, 64, 81, 125, 343, 729
- $size$: 25, 30, 35, 40, 45
- Taint run: $p = 8$, $size = 5$
- Taint overhead: < 5 mins

Case Studies

Piz Daint, 21 nodes

- Intel Xeon E5-2695 v4 2.1 GHz
- 2 sockets, 18 cores each
- 128 GB Memory
- Cray MPICH 7.7.2

Skylake Cluster, 2 nodes

- Intel Xeon 6154 3 GHz
- 36 cores
- 384 GB Memory
- OpenMPI 4.0.3

LULESH

- p : 27, 64, 81, 125, 343, 729
- $size$: 25, 30, 35, 40, 45
- Taint run: $p = 8$, $size = 5$
- Taint overhead: < 5 mins


MILC su3_rmd

- p : 4, 8, 16, 32, 64
- $size$: 32, 64, 128, 256, 512
- Taint run: $p = 32$, $size = 128$
- Taint overhead: ~1 hr


Parameter Pruning

LULESH	Total	Comm	p	size	regions	iters	balance	cost		p, size
Functions [count]	349	2 + 7	2	40	13	4	9	2		40
Loops [count]	275	-	2	78	27	4	20	2		78
MILC	Total	Comm	p	size	trajecs	warms steps	nrest. niter	mass, beta, nfl.	u0	p, size
Functions [count]	629	13 + 8	54	53	12	9	6	1	4	56
Loops [count]	874	-	187	161	39	31	15	1	7	196




Instrumentation Overhead




Instrumentation Overhead






Instrumentation Overhead






Instrumentation Overhead






Instrumentation Overhead



Instrumentation Overhead



Instrumentation Overhead



Less Noise

```

int foo(int a, int b, int& result) {
    for(int i = 0 ; i < a; ++i)
        result += b * i;
}
    
```

$$0.5a + 10^{-3}b$$

Separate **program**
from **noise**.

Less Noise

```
int foo(int a, int b, int& result) {
    for(int i = 0 ; i < a; ++i)
        result += b * i;
}
```

$$0.5a + 10^{-3}b$$

Separate **program**
from **noise**.

Less Noise

```
int foo(int a, int b, int& result) {
    for(int i = 0 ; i < a; ++i)
        result += b * i;
}
```

$$0.5a + 10^{-3}b$$

Separate program
from noise.

Less Noise

```
int foo(int a, int b, int& result) {  
    for(int i = 0 ; i < a; ++i)  
        result += b * i;  
}
```

$$0.5a + 10^{-3}b$$

Separate **program**
from **noise**.

LULESH

- 86.2% of functions are constant
- TOP 5 models with perf-taint: parse 14 functions
- Same TOP 5 models with black-box modeling: parse 33 functions

Less Noise

```
int foo(int a, int b, int& result) {  
    for(int i = 0 ; i < a; ++i)  
        result += b * i;  
}
```

$$0.5a + 10^{-3}b$$

Separate **program**
from **noise**.

LULESH

- 86.2% of functions are constant
- TOP 5 models with perf-taint: parse 14 functions
- Same TOP 5 models with black-box modeling: parse 33 functions

MILC su3_rmd

- 87.7% of functions are constant
- TOP 5 models with perf-taint: parse 32 functions
- Same TOP 5 models with black-box modeling: parse 43 functions

Less Intrusion

```
int bar(int a) {
    instrument();
    return a * a;
}
```

```
int foo(int a, int& res) {
    instrument();
    for(int i = 0 ; i < a; ++i)
        res += bar(i);
}
```

$$1.3a + 10^{-4} \sqrt{a}$$

Separate **program**
and **instrumentation**.

Less Intrusion

```
int bar(int a) {
  instrument();
  return a * a;
}
```

```
int foo(int a, int& res) {
  instrument();
  for(int i = 0 ; i < a; ++i)
    res += bar(i);
}
```

$$1.3a + 10^{-4} \sqrt{a}$$

Separate **program**
and **instrumentation**.

Less Intrusion

```
int bar(int a) {
    instrument();
    return a * a;
}
```

```
int foo(int a, int& res) {
    instrument();
    for(int i = 0 ; i < a; ++i)
        res += bar(i);
}
```

$$1.3a + 10^{-4} \sqrt{a}$$

Separate **program**
and **instrumentation**.

Less Intrusion

```
int bar(int a) {  
  instrument();  
  return a * a;  
}
```

```
int foo(int a, int& res) {  
  instrument();  
  for(int i = 0 ; i < a; ++i)  
    res += bar(i);  
}
```

$$1.3a + 10^{-4} \sqrt{a}$$

Separate **program**
and **instrumentation**.



$$3 * 10^{-3} * p^{0.5} + 10^{-5} * size^3$$

Less Intrusion

```
int bar(int a) {  
  instrument();  
  return a * a;  
}
```

```
int foo(int a, int& res) {  
  instrument();  
  for(int i = 0 ; i < a; ++i)  
    res += bar(i);  
}
```

$$1.3a + 10^{-4} \sqrt{a}$$

Separate **program**
and **instrumentation**.




$$3 * 10^{-3} * p^{0.5} + 10^{-5} * size^3$$




$$2.4 * 10^{-8} * p^{0.25} * size^3$$


Hardware Contention




Hardware Contention




Hardware Contention




Hardware Contention



Hardware Contention



Hardware Contention



Experiment Design

```
int foo(int a) {
    if(a < 4) kernel_linear(a);
    else     kernel_log(a);
}
```

$$f: \begin{cases} a & a < 4 \\ \log_2 a & a \geq 8 \end{cases}$$

Qualitative change of behavior.
 Modeling requires **consistency**.

Experiment Design


```
int foo(int a) {
    if(a < 4) kernel_linear(a);
    else     kernel_log(a);
}
```

$$f: \begin{cases} a & a < 4 \\ \log_2 a & a \geq 8 \end{cases}$$

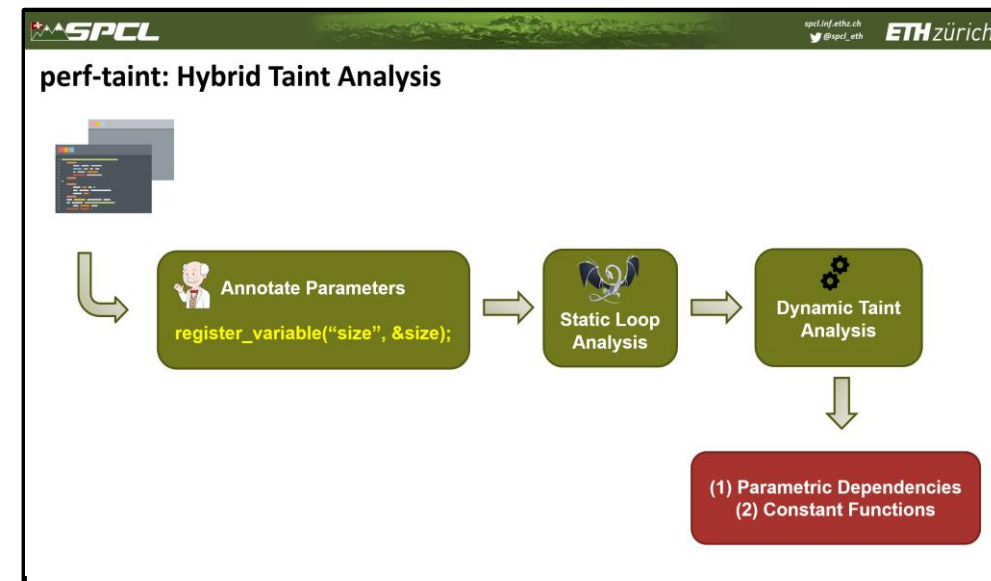
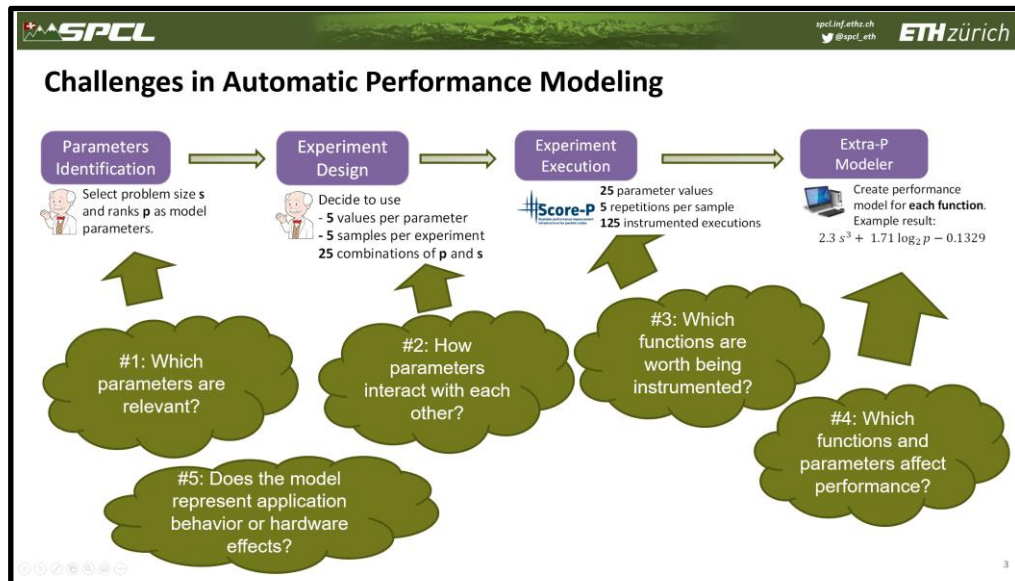
Qualitative change of behavior.
Modeling requires **consistency**.

Summary

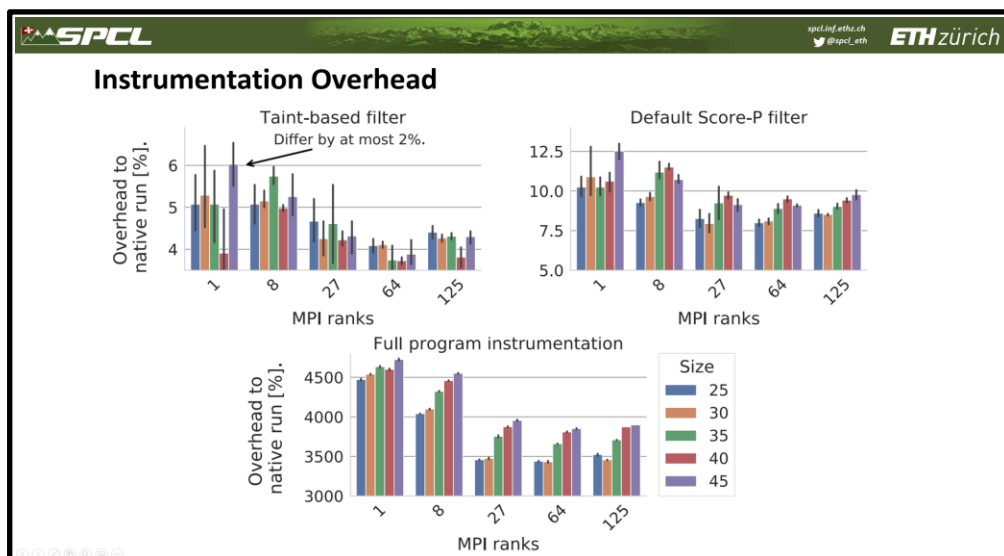
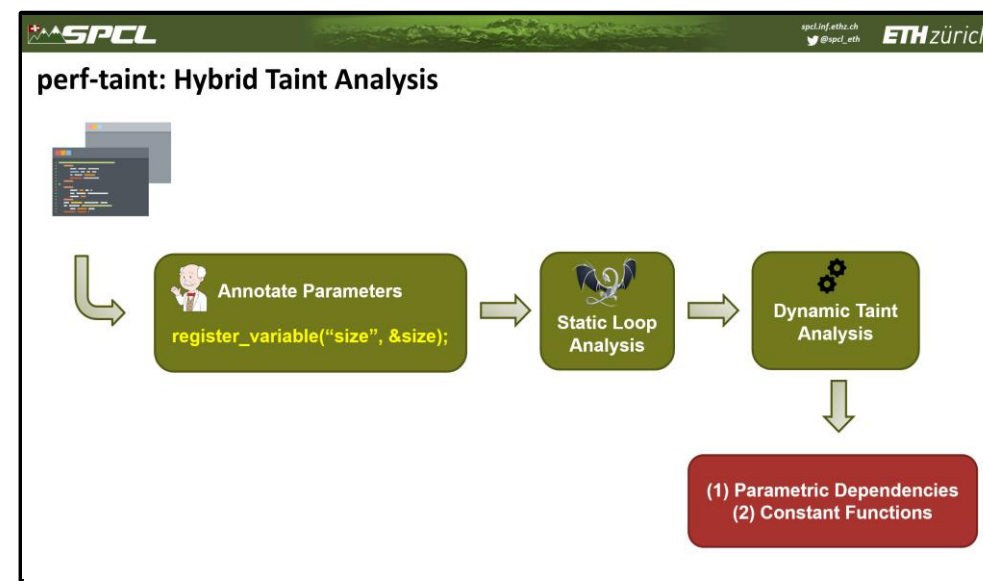
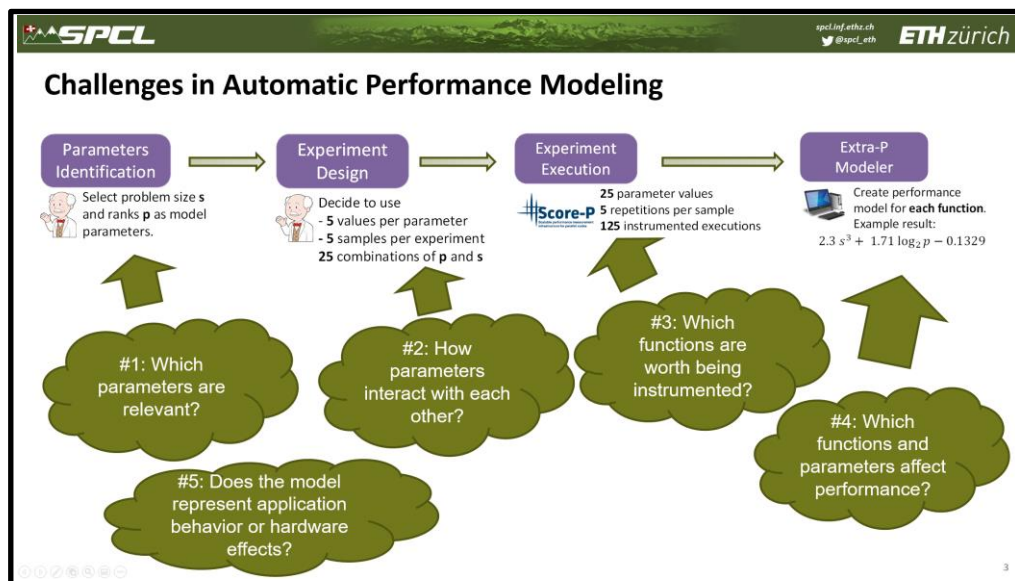
Summary




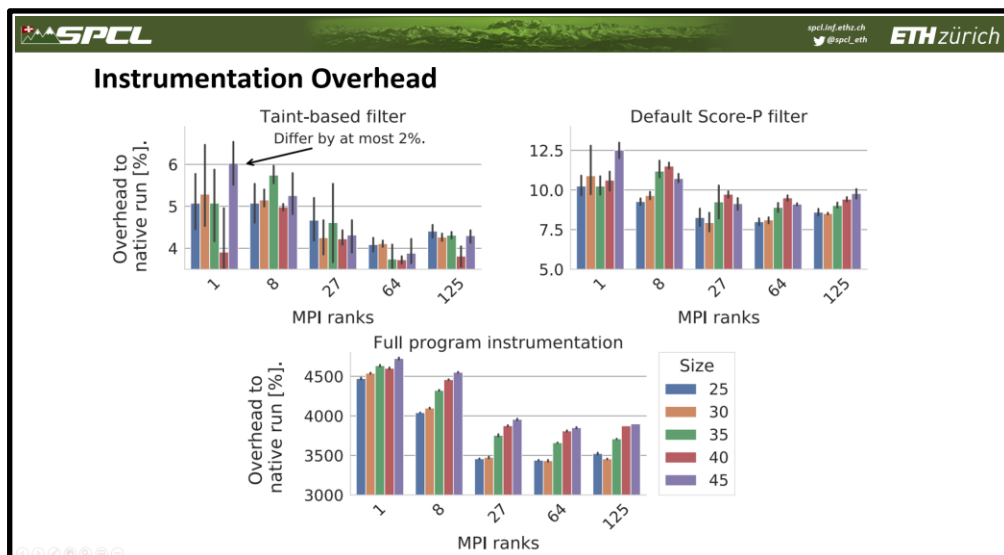
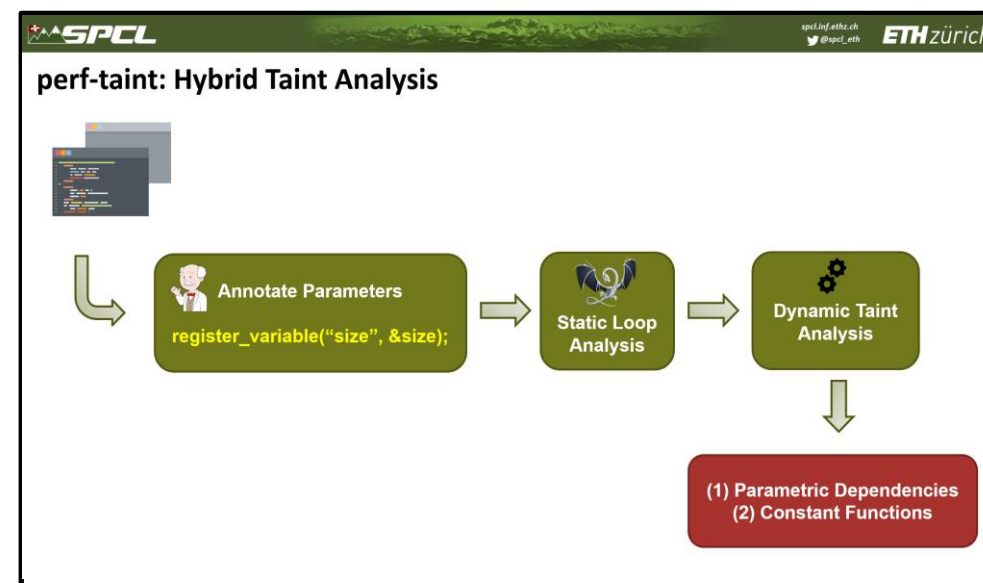
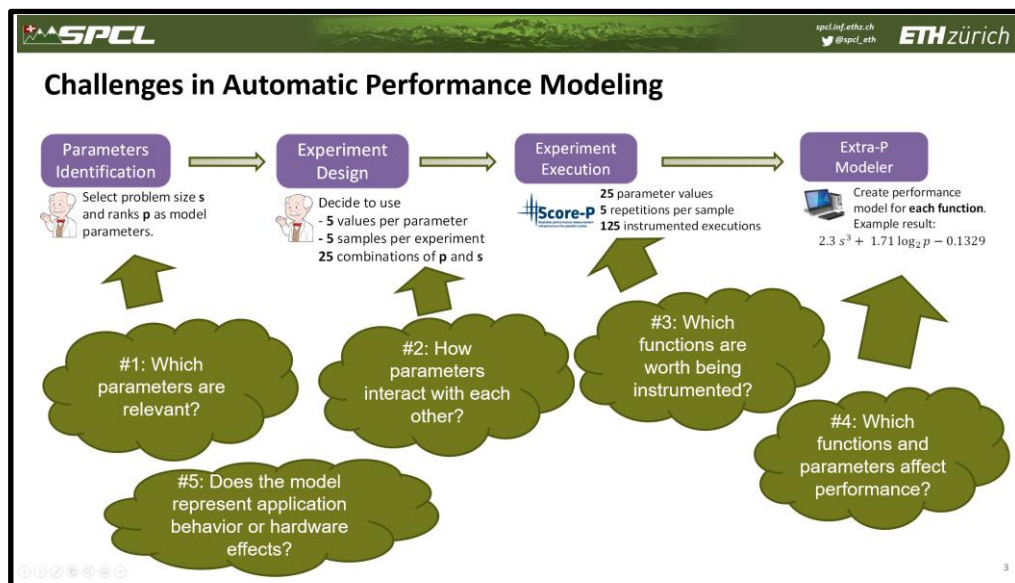
Summary



Summary



Summary



spcl/perf-taint